

Query processing and optimization

Definitions

- Query processing
 - translation of query into low-level activities
 - evaluation of query
 - data extraction
- Query optimization
 - selecting the most efficient query evaluation

Query Processing (1/2)

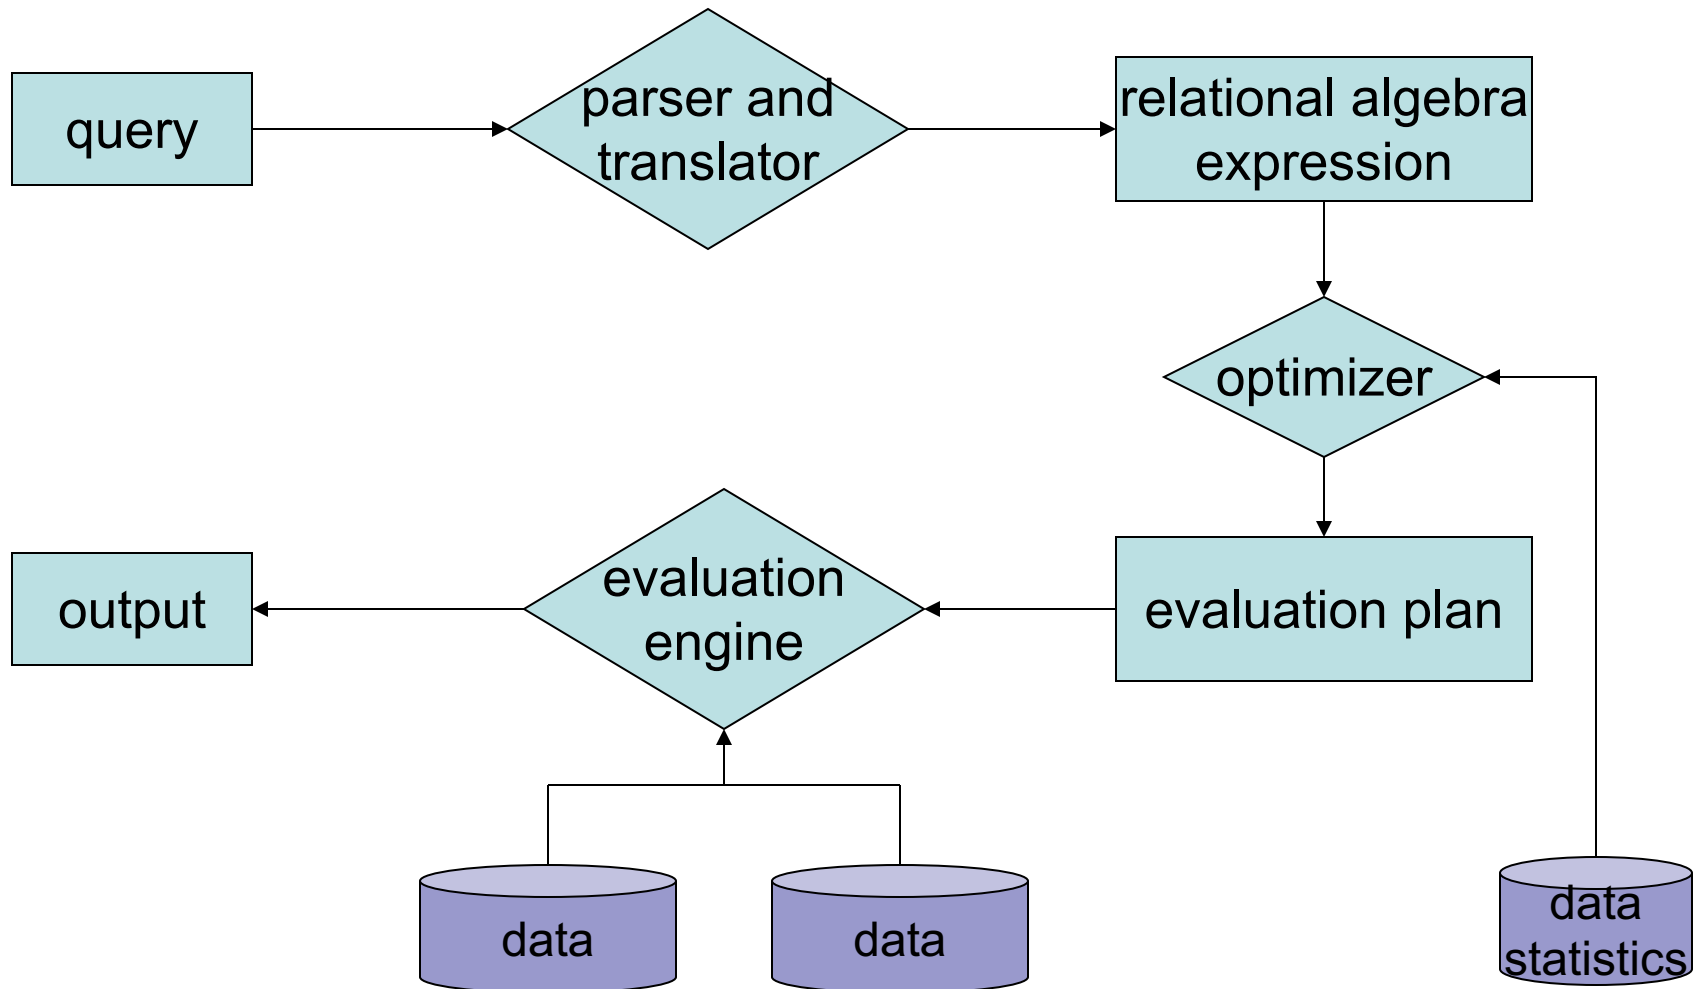
- SELECT * FROM student WHERE name=Paul
- Parse query and translate
 - check syntax, verify names, etc
 - translate into relational algebra (RDBMS)
 - create **evaluation plans**
- Find best plan (optimization)
- Execute plan

student	
<u>cid</u>	name
00112233	Paul
00112238	Rob
00112235	Matt

takes	
<u>cid</u>	<u>courseid</u>
00112233	312
00112233	395
00112235	312

course	
<u>courseid</u>	<u>coursename</u>
312	Advanced DBs
395	Machine Learning

Query Processing (2/2)



Relational Algebra (1/2)

- Query language
- Operations:
 - select: σ
 - project: π
 - union: \cup
 - difference: $-$
 - product: \times
 - join: \bowtie
- Extended relational algebra operations:

Relational Algebra (2/2)

- SELECT * FROM student WHERE name=Paul
 - $\sigma_{\text{name=Paul}}(\text{student})$
- $\pi_{\text{name}}(\sigma_{\text{cid}<00112235}(\text{student}))$
- $\pi_{\text{name}}(\sigma_{\text{coursename=Advanced DBs}}((\text{student} \bowtie_{\text{cid}} \text{takes}) \bowtie_{\text{courseid}} \text{course}))$

student	
<u>cid</u>	name
00112233	Paul
00112238	Rob
00112235	Matt

takes	
<u>cid</u>	<u>courseid</u>
00112233	312
00112233	395
00112235	312

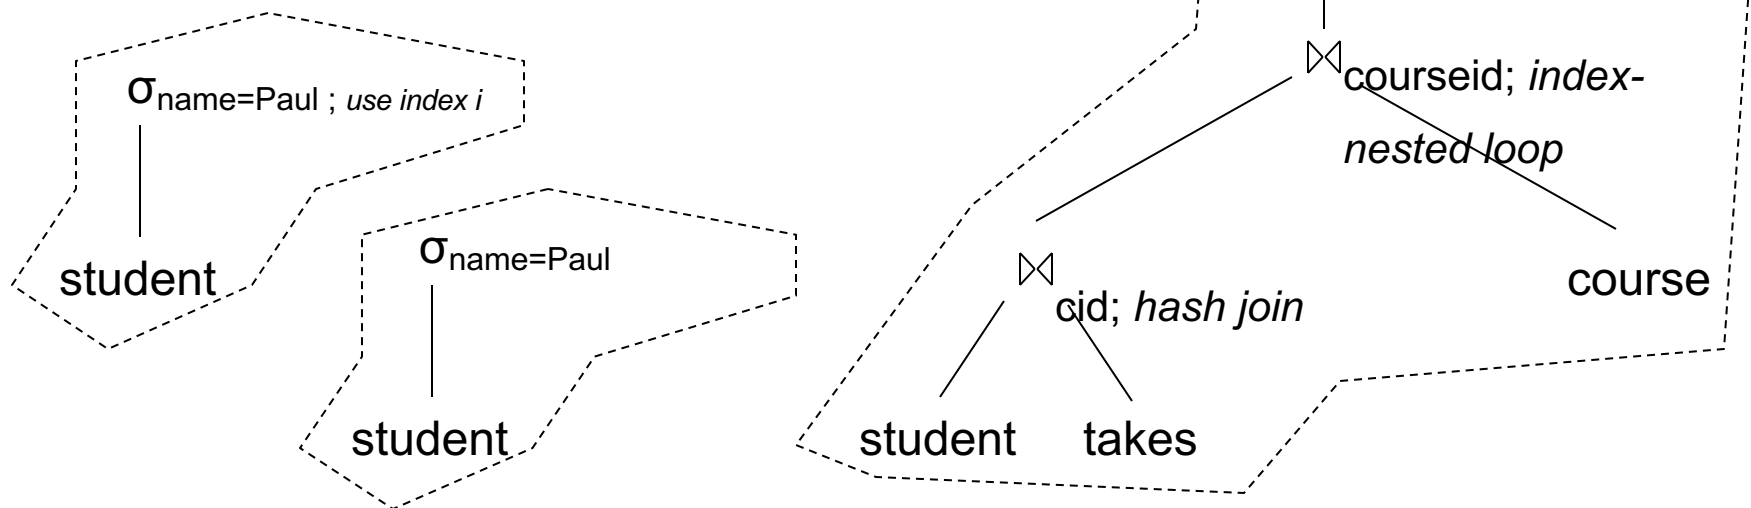
course	
<u>courseid</u>	<u>coursename</u>
312	Advanced DBs
395	Machine Learning

Why Optimize?

- Many alternative options to evaluate a query
 - $\pi_{\text{name}}(\sigma_{\text{course name}=\text{Advanced DBs}}((\text{student} \bowtie_{\text{cid}} \text{takes}) \bowtie_{\text{courseid}} \text{course}))$
 - $\pi_{\text{name}}((\text{student} \bowtie_{\text{cid}} \text{takes}) \bowtie_{\text{courseid}} \sigma_{\text{course name}=\text{Advanced DBs}}(\text{course}))$
- Several options to evaluate a single operation
 - $\sigma_{\text{name}=\text{Paul}}(\text{student})$
 - scan file
 - use secondary index on student.name
- Multiple access paths
 - access path: how can records be accessed (using index; which index?; etc.)

Evaluation plans

- Specify which access path to follow
- Specify which algorithm to use to evaluate operator
- Specify how operators interleave
- Optimization:
 - estimate the cost of each plan (not all plans)
 - select plan with lowest estimated cost



Estimating Cost

- What needs to be considered:
 - Disk I/Os
 - sequential (reading neighbouring pages is faster)
 - random
 - CPU time
 - Network communication
- What are we going to consider:
 - Disk I/Os
 - page (data block) reads/writes
 - Ignoring cost of writing final output

Operations and Costs (1/2)

- Operations: σ , π , \cup , \cap , $-$, \times , \bowtie
- Costs:
 - N_R : number of records in R (other notation: T_R or $T(R) \rightarrow$ tuple)
 - L_R : size of record in R (length of record)
 - bf_R : blocking factor (other notation: F_R)
 - number of records in a page (datablock)
 - B_R : number of pages to store relation R
 - $V(R, A)$: number of distinct values of attribute A in R
other notation: $I_A(R)$ (Image size)
 - $SC(R, A)$: selection cardinality of A in R (number of matching rec.)
 - A key: $SC(R, A)=1$
 - A nonkey: $SC(R, A)= T_R / V(R, A)$ (uniform distribution assumption)
 - HT_i : number of levels in index I (\rightarrow height of tree)
 - rounding up fractions and logarithms

Operations and Costs (2/2)

- relation *takes*
 - 7000 tuples
 - student cid 8 bytes
 - course id 4 bytes
 - 40 courses
 - 1000 students
 - page size 512 bytes
 - output size (in pages) of query:
which students take the Advanced DBs course?
 - $T_{\text{takes}} = 7000$
 - $V(\text{courseid}, \text{takes}) = 40$
 - $SC(\text{courseid}, \text{takes}) = \text{ceil}(T_{\text{takes}}/V(\text{courseid}, \text{takes})) = \text{ceil}(7000/40) = 175$
 - $bf_{\text{takes}} = \text{floor}(512/12) = 42$ $B_{\text{takes}} = 7000/42 = 167$ pages
 - $bf_{\text{output}} = \text{floor}(512/8) = 64$ $B_{\text{output}} = 175/64 = 3$ pages

Cost of Selection σ (1/2)

- Linear search
 - read all pages, find records that match (assuming equality search)
 - average cost:
 - nonkey (multiple occurrences): B_R , key: $0.5 \cdot B_R$
- Binary search
 - on ordered field
 - average cost: $\lceil \log_2 B_R \rceil + m$
 - m additional pages to be read (first found then read the duplicates)
 - $m = \text{ceil}(SC(A,R)/bf_R) - 1$
- Primary/Clustered Index (B+ tree)
 - average cost:
 - single record: $HT_i + 1$
 - multiple records: $HT_i + \text{ceil}(SC(R,A)/bf_R)$

Cost of Selection σ (2/2)

- Secondary Index (B+ tree)
 - average cost:
 - key field: $HT_i + 1$
 - nonkey field
 - worst case $HT_i + SC(A,R)$
 - linear search more desirable if many matching records !!!

Complex selection σ_{expr}

- conjunctive selections: $\sigma_{\theta_1 \wedge \theta_2 \dots \wedge \theta_n}$
 - perform simple selection using θ_i with the lowest evaluation cost
 - e.g. using an index corresponding to θ_i
 - apply remaining conditions θ on the resulting records
 - $\sigma_{cid > 00112233 \wedge courseid = 312}(\text{takes})$
 - cost: the cost of the simple selection on selected θ
 - multiple indices
 - select indices that correspond to θ_i s
 - scan indices and return RIDs (ROWID in Oracle)
 - answer: intersection of RIDs
 - cost: the sum of costs + record retrieval
- disjunctive selections: $\sigma_{\theta_1 \vee \theta_2 \dots \vee \theta_n}$
 - multiple indices
 - union of RIDs
 - linear search

Projection and set operations

- SELECT DISTINCT cid FROM takes
 - π requires duplicate elimination
 - sorting
- set operations require duplicate elimination
 - $R \cap S$
 - $R \cup S$
 - sorting

Sorting

- efficient evaluation for many operations
- required by query:
 - `SELECT cid, name FROM student ORDER BY name`
- implementations
 - **internal sorting** (if records fit in memory)
 - **external sorting**
(that's why we need temporary space on disk)

External Sort-Merge Algorithm (1/3)

- Sort stage: create sorted *runs*

i=0;

repeat

 read M pages of relation R into memory (**M : size of Memory**)

 sort the M pages

 write them into file R_i

 increment i

until no more pages

$N = i$ // number of runs

External Sort-Merge Algorithm (2/3)

- Merge stage: merge sorted *runs*

//assuming $N < M$ ($N \leq M-1$ we need 1 output buffer)

allocate a page for each run file R_i // N pages allocated

read a page P_i of each R_i

repeat

 choose first record (in sort order) among N pages, say from page P_j

 write record to output and delete from page P_j


 if page is empty read next page P_j' from R_j

until all pages are empty

External Sort-Merge Algorithm (3/3)

- Merge stage: merge sorted *runs*
- What if $N \geq M$?
 - perform multiple *passes*
 - each *pass* merges $M-1$ runs until relation is processed
 - in next pass number of runs is reduced
 - final *pass* generated sorted output

Sort-Merge Example



Sort-Merge cost

- B_R the number of pages of R
- Sort stage: $2 * B_R$
 - read/write relation
- Merge stage:
 - initially $\left\lceil \frac{B_R}{M} \right\rceil$ runs to be merged
 - each pass $M-1$ runs sorted
 - thus, total number of passes: $\left\lceil \log_{M-1} \left(\frac{B_R}{M} \right) \right\rceil$
 - at each pass $2 * B_R$ pages are read/written
 - read/write relation ($B_R + B_R$)
 - apart from final write (B_R)
- Total cost:
 - $2 * B_R + 2 * B_R * \left\lceil \log_{M-1} \left(\frac{B_R}{M} \right) \right\rceil - B_R$ eg. $B_R = 1000000, M=100$

Projection

- $\pi_{A_1, A_2 \dots} (R)$
- remove unwanted attributes
 - scan and drop attributes
- remove duplicate records
 - sort resulting records using all attributes as sort order
 - scan sorted result, eliminate duplicates (adjacent)
- cost
 - initial scan + **sorting** + final scan

Join

- $\pi_{\text{name}}(\sigma_{\text{course name}=\text{Advanced DBs}}((\text{student} \bowtie_{\text{cid}} \text{takes}) \bowtie_{\text{courseid}} \text{course}))$
- implementations
 - nested loop join
 - block-nested loop join
 - indexed nested loop join
 - sort-merge join
 - hash join

Nested loop join (1/2)

- $R \bowtie S$

for each tuple t_R of R

 for each t_S of S

 if (t_R t_S match) output $t_R.t_S$

 end

end

- Works for any join condition
- S inner relation
- R outer relation

Nested loop join (2/2)

- Costs:
 - **best case** when smaller relation fits in memory
 - use it as inner relation
 - $B_R + B_S$
 - **worst case** when memory holds one page of each relation
 - S scanned for each tuple in R
 - $T_R * B_S + B_R$

Block nested loop join (1/2)

```
for each page  $X_R$  of R
  for each page  $X_S$  of S
    for each tuple  $t_R$  in  $X_R$ 
      for each  $t_S$  in  $X_S$ 
        if ( $t_R$   $t_S$  match) output  $t_R.t_S$ 
      end
    end
  end
end
end
```

Block nested loop join (2/2)

- Costs:
 - **best case** when smaller relation fits in memory
 - use it as inner relation
 - $B_R + B_S$
 - **worst case** when memory holds one page of each relation
 - S scanned for each page in R
 - $B_R * B_S + B_R$

Block nested loop join (an improvement)

Memory size: M

for each $M - 1$ page size chunk M_R of R

for each page X_S of S

for each tuple t_R in M_R

for each t_S in X_S

if (t_R t_S match) output $t_R.t_S$

end

end

end

end

Block nested loop join (an improvement)

- Costs:
 - best case when smaller relation fits in memory
 - use it as inner relation
 - $B_R + B_S$
 - general case
 - S scanned for each M-1 size chunk in R
 - $(B_R / (M-1)) * B_S + B_R$

Indexed nested loop join

- $R \bowtie S$
- Index on inner relation (S)
- for each tuple in outer relation (R) *probe* index of inner relation
- Costs:
 - $B_R + T_R * c$
 - c the cost of index-based selection of inner relation
 - $c \approx T(s)/V(s,A)$
(if A is the join column and index is kept in memory)
 - relation with fewer records as outer relation


Sort-merge join

- $R \bowtie S$
- Relations sorted on the join attribute
- Merge sorted relations
 - pointers to first record in each relation
 - read in a group of records of S with the same values in the join attribute
 - read records of R and process
- Relations in sorted order to be read once
- Cost:
 - cost of sorting + $B_S + B_R$

→	d	D	→	e	67
	e	E		e	87
	x	X		n	11
	v	V		v	22
				z	38

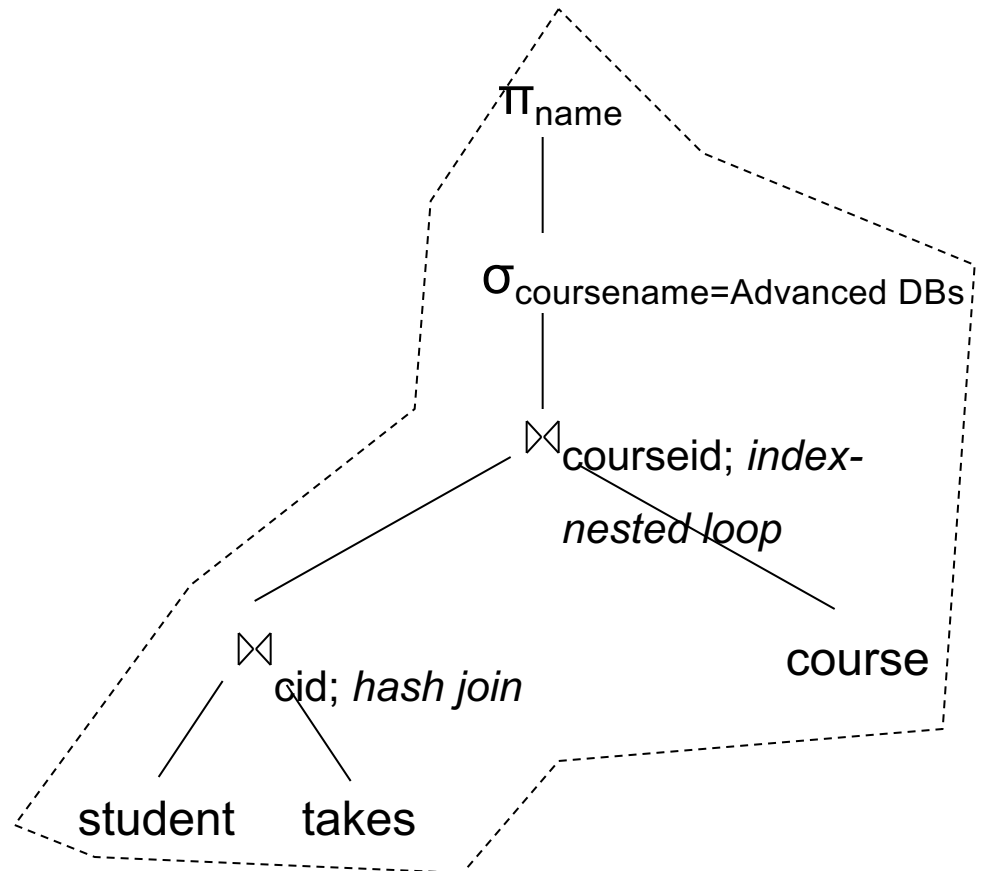
Hash join

- $R \bowtie S$
- use h_1 on joining attribute to map records to partitions that fit in memory
 - records of R are partitioned into $R_0 \dots R_{n-1}$
 - records of S are partitioned into $S_0 \dots S_{n-1}$
- join records in corresponding partitions
 - using a hash-based indexed block nested loop join
- Cost: $2 \cdot (B_R + B_S) + (B_R + B_S)$



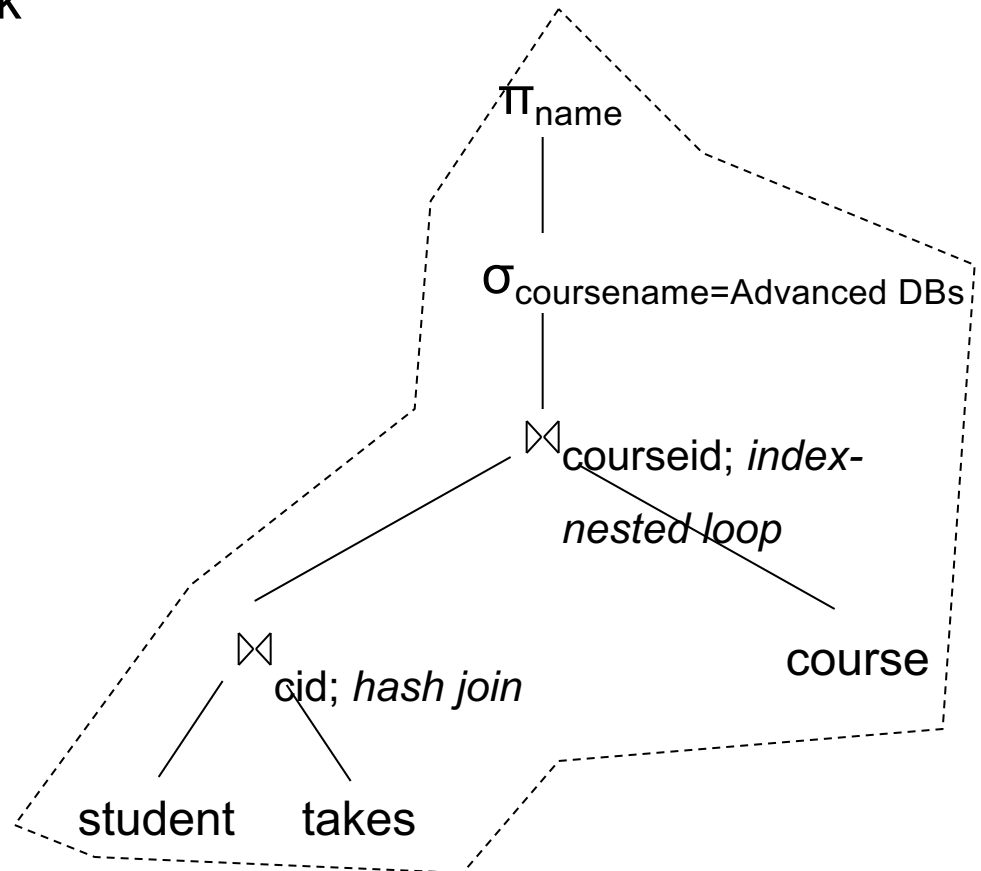
Evaluation

- evaluate multiple operations in a plan
- materialization
- pipelining



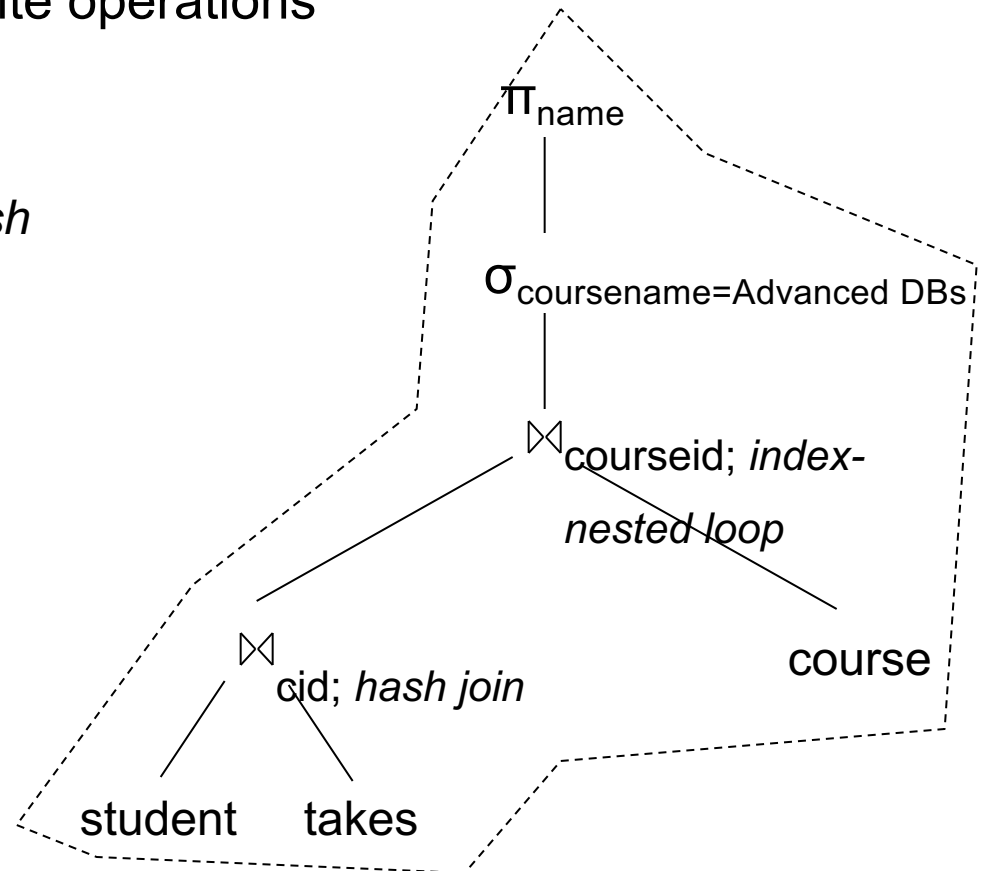
Materialization

- create and read temporary relations
- create implies writing to disk
 - more page writes



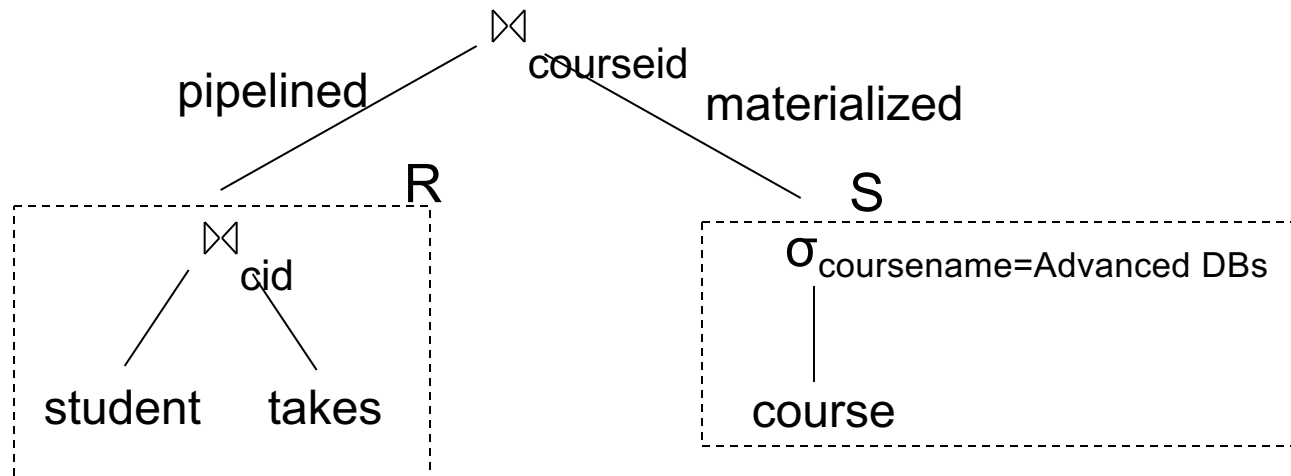
Pipelining (1/2)

- creating a pipeline of operations
- reduces number of read-write operations
- implementations
 - demand-driven - data *pull*
 - producer-driven - data *push*



Pipelining (2/2)

- can pipelining always be used?
- any algorithm?
- cost of $R \bowtie S$
 - materialization and hash join: $B_R + 3(B_R + B_S)$
 - pipelining and indexed nested loop join: $T_R * HT_i$

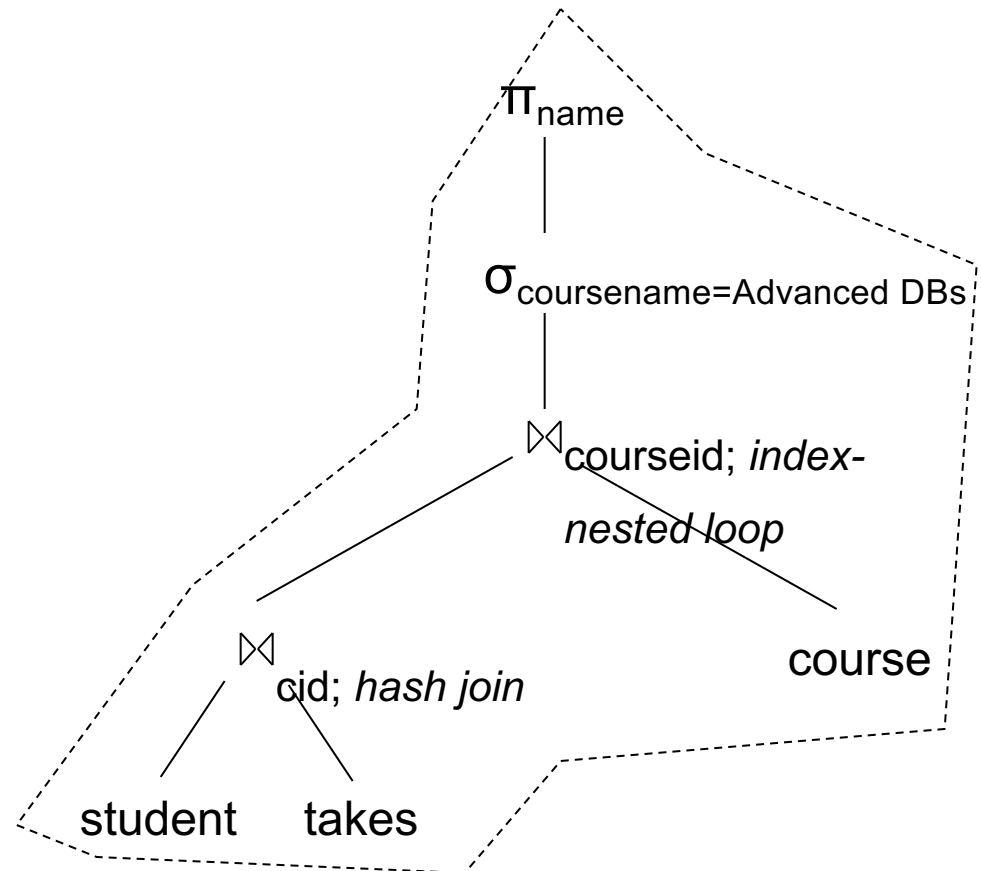


Choosing evaluation plans

- cost based optimization
- enumeration of plans
 - $R \bowtie S \bowtie T$, 12 possible orders ($3! * 2$)
 $(R \bowtie S) \bowtie T$, $R \bowtie (S \bowtie T)$
- cost estimation of each plan
- overall cost
 - cannot optimize operation independently

Cost estimation

- operation (σ , π , \bowtie ...)
- implementation
- size of inputs
- size of outputs
- sorting




Expression Equivalence

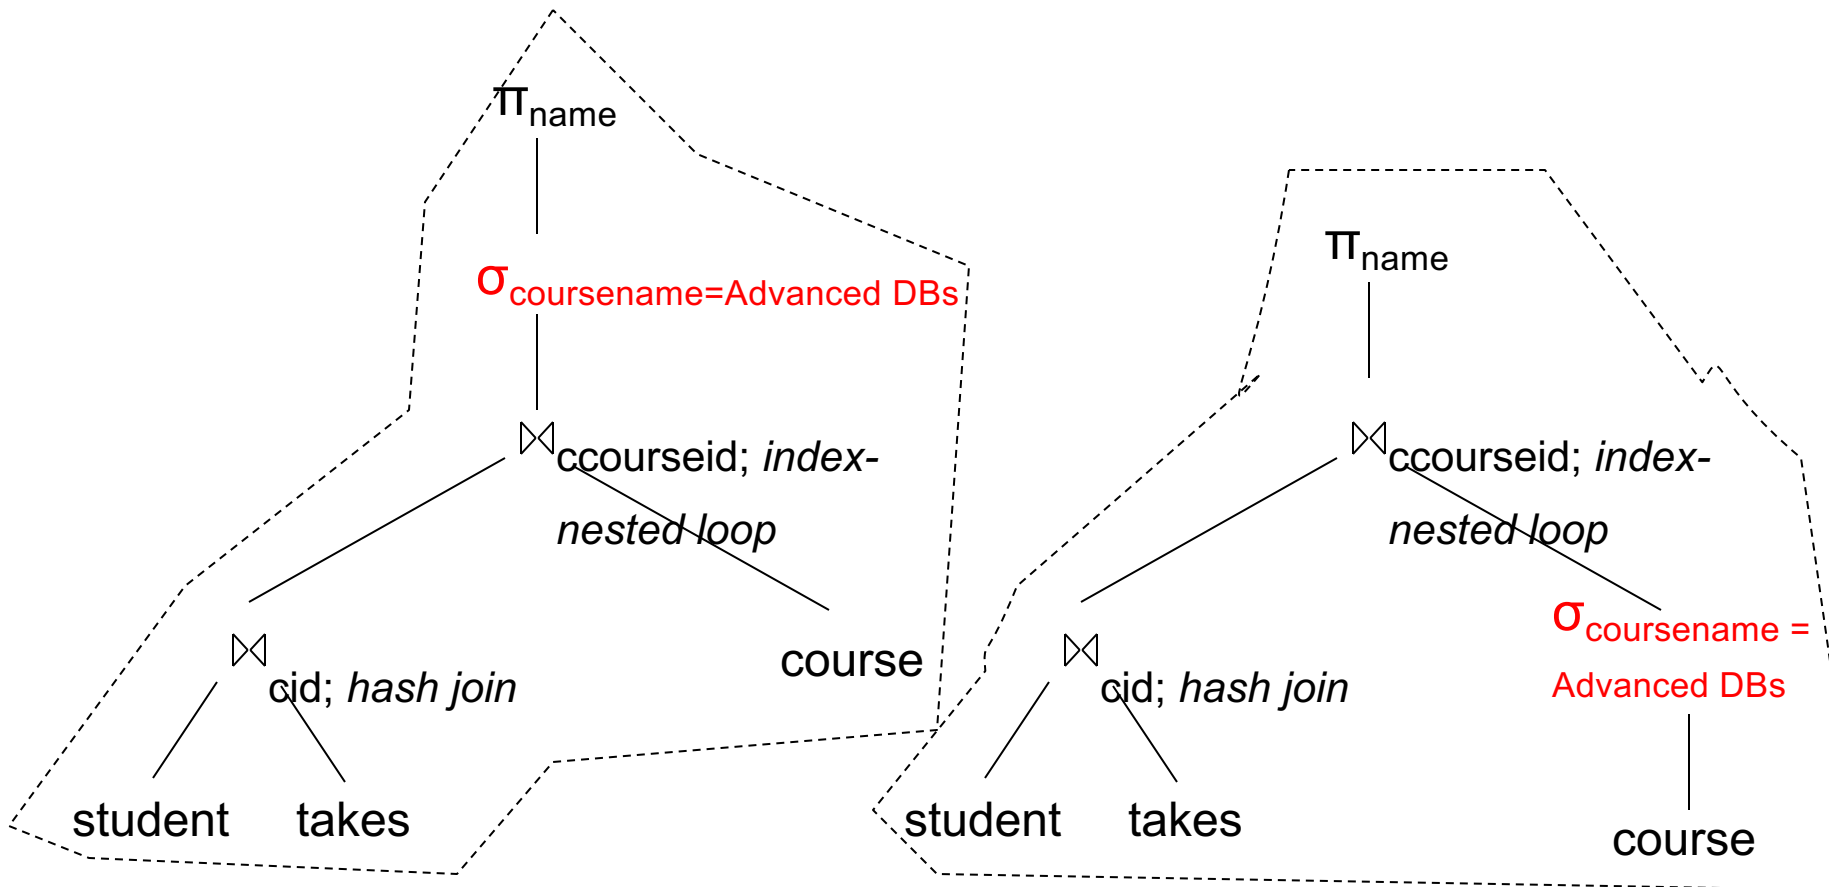
- conjunctive selection decomposition
 - $\sigma_{\theta_1 \wedge \theta_2}(R) = \sigma_{\theta_1}(\sigma_{\theta_2}(R))$
- commutativity of selection
 - $\sigma_{\theta_1}(\sigma_{\theta_2}(R)) = \sigma_{\theta_2}(\sigma_{\theta_1}(R))$
- combining selection with join and product
 - $\sigma_{\theta_1}(R \bowtie S) = \dots \quad R \bowtie_{\theta_1} S = \dots$
- commutativity of joins
 - $R \bowtie_{\theta_1} S = S \bowtie_{\theta_1} R$
- distribution of selection over join
 - $\sigma_{\theta_1 \wedge \theta_2}(R \bowtie S) = \sigma_{\theta_1}(R) \bowtie \sigma_{\theta_2}(S)$
- distribution of projection over join
 - $\pi_{A_1, A_2}(R \bowtie S) = \pi_{A_1}(R) \bowtie \pi_{A_2}(S)$
- associativity of joins: $R \bowtie (S \bowtie T) = (R \bowtie S) \bowtie T$

Cost Optimizer (1/2)

- transforms expressions
 - equivalent expressions
 - heuristics, *rules of thumb*
 - perform **selections early**
 - perform **projections early**
 - replace products followed by selection σ ($R \times S$) with joins $R \bowtie S$
 - start with joins, selections with smallest result
 - create **left-deep join trees**



Cost Optimizer (2/2)



Summary

- Estimating the cost of a single operation
- Estimating the cost of a query plan
- Optimization
 - choose the most efficient plan