


# Notes 09: Concurrency Control

Ullman et al. :  
Database System Principles

# Chapter 18 [18] Concurrency Control



Interactions among concurrently executing transactions can cause **the database state to become inconsistent, even when the transactions individually preserve correctness of the state, and there is no system failure.**

Thus, the timing of individual steps of different transactions needs to be regulated in some manner.

This regulation is the job of the *scheduler component of the DBMS*, and **the general process of assuring that transactions preserve consistency when executing simultaneously** is called ***concurrency control***.

In most situations, the **scheduler will execute the reads and writes directly**, first calling on the buffer manager if the desired database element is not in a buffer.

However, in some situations, it is not safe for the request to be executed immediately.

The **scheduler must delay the request**.

In some concurrency-control techniques, the scheduler may even abort the transaction that issued the request

## Example:

T1: Read(A)  
A  $\leftarrow$  A+100  
Write(A)  
Read(B)  
B  $\leftarrow$  B+100  
Write(B)

T2: Read(A)  
A  $\leftarrow$  A $\times$ 2  
Write(A)  
Read(B)  
B  $\leftarrow$  B $\times$ 2  
Write(B)

Constraint: A=B

A ***schedule*** is a sequence of the important actions taken by one or more transactions.

When studying concurrency control, the important read and write actions take place in the main-memory buffers, not the disk.

That is, a database element  $A$  that is brought to a buffer by some transaction  $T$  may be read or written in that buffer not only by  $T$  but by other transactions that access  $A$ .

A ***schedule*** is ***serial*** if its actions consist of all the actions of one transaction, then all the actions of another transaction, and so on.

No mixing of the actions is allowed.

# Schedule A (serial)

T1	T2	A	B
		25	25
Read(A); $A \leftarrow A + 100$			
Write(A);		125	
Read(B); $B \leftarrow B + 100$ ;			
Write(B);			125
	Read(A); $A \leftarrow A \times 2$ ;		
	Write(A);	250	
	Read(B); $B \leftarrow B \times 2$ ;		
	Write(B);		250
		250	250

# Schedule B (serial)

T1	T2	A	B
		25	25
	Read(A); $A \leftarrow A \times 2$ ;		
	Write(A);	50	
	Read(B); $B \leftarrow B \times 2$ ;		
	Write(B);		50
Read(A); $A \leftarrow A + 100$			
Write(A);		150	
Read(B); $B \leftarrow B + 100$ ;			
Write(B);			150
		150	150



A **schedule  $S$  is serializable** if there is a serial schedule  $S'$  such that for every initial database state, the **effects** of  $S$  and  $S'$  are the **same**.

# Schedule C (serializable)

T1	T2	A	B
		25	25
Read(A); $A \leftarrow A + 100$			
Write(A);		125	
	Read(A); $A \leftarrow A \times 2$ ;		
	Write(A);	250	
Read(B); $B \leftarrow B + 100$ ;			125
Write(B);			
	Read(B); $B \leftarrow B \times 2$ ;		
	Write(B);		250
		250	250

# Schedule D (not serializable)

T1	T2	A	B
Read(A); $A \leftarrow A+100$		25	25
Write(A);		125	
	Read(A); $A \leftarrow A \times 2$ ;		
	Write(A);	250	
	Read(B); $B \leftarrow B \times 2$ ;		50
	Write(B);		150
Read(B); $B \leftarrow B+100$ ;			
Write(B);		250	150

# Schedule E

Same as Schedule D  
but with new T2'

T1	T2'	A	B
		25	25
Read(A); $A \leftarrow A + 100$			
Write(A);		125	
	Read(A); $A \leftarrow A \times 1$ ;		
	Write(A);	125	
	Read(B); $B \leftarrow B \times 1$ ;		
	Write(B);		25
Read(B); $B \leftarrow B + 100$ ;			
Write(B);			125
		125	125

- Want schedules that are “good”, regardless of
  - initial state and
  - transaction semantics
- Only look at **order of read and writes**

It is not realistic for the scheduler to concern itself with the details of computation undertaken by transactions.


Any database element  $A$  *that a transaction  $T$  writes is given a value* that depends on the database state in such a way that **no arithmetic coincidences occur**.

To make the notation precise:

1. An *action* is an expression of the form  $r_i(X)$  or  $w_i(X)$ , meaning that transaction  $T_i$ , *reads* or *writes*, respectively, the database element  $X$ .
2. A *transaction*  $T_i$  is a sequence of actions with subscript  $i$ .
3. A *schedule*  $S$  of a set of transactions  $T$  is a *sequence of actions*, in which for each transaction  $T_i$  in  $T$ , the actions of  $T_i$  appear in  $S$  *in the same order* that they appear in the definition of  $T_i$  itself. We say that  $S$  is an *interleaving of the actions of the transactions of which it is composed*.

Example:

$Sc = r_1(A)w_1(A)r_2(A)w_2(A)r_1(B)w_1(B)r_2(B)w_2(B)$



$Sc' = r_1(A)w_1(A) r_1(B)w_1(B)r_2(A)w_2(A)r_2(B)w_2(B)$

$T_1$

$T_2$

$Sc$  is "equivalent" to a serial schedule.

$Sc$  is "good"


However, for Sd:

$S_d = \underbrace{r_1(A)w_1(A)}_{\text{}} \underbrace{r_2(A)w_2(A)}_{\text{}} \underbrace{r_2(B)w_2(B)}_{\text{}} \underbrace{r_1(B)w_1(B)}_{\text{}}$

- as a matter of fact,  
     $T_2$  must precede  $T_1$   
    in any equivalent schedule,  
    i.e.,  $T_2 \rightarrow T_1$



- $T_2 \rightarrow T_1$
- Also,  $T_1 \rightarrow T_2$



$S_d$  cannot be rearranged  
into a serial schedule




$S_d$  is not “equivalent” to  
any serial schedule



$S_d$  is “bad”

## Returning to Sc

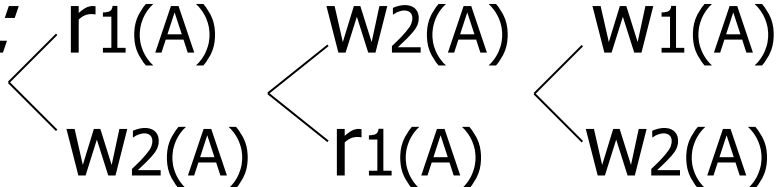


☛ **no cycles**  $\Rightarrow$  Sc is “equivalent” to a serial schedule  
(in this case  $T_1, T_2$ )

# Concepts

*Transaction:* sequence of  $r_i(x)$ ,  $w_i(x)$  actions

*Conflicting actions:*

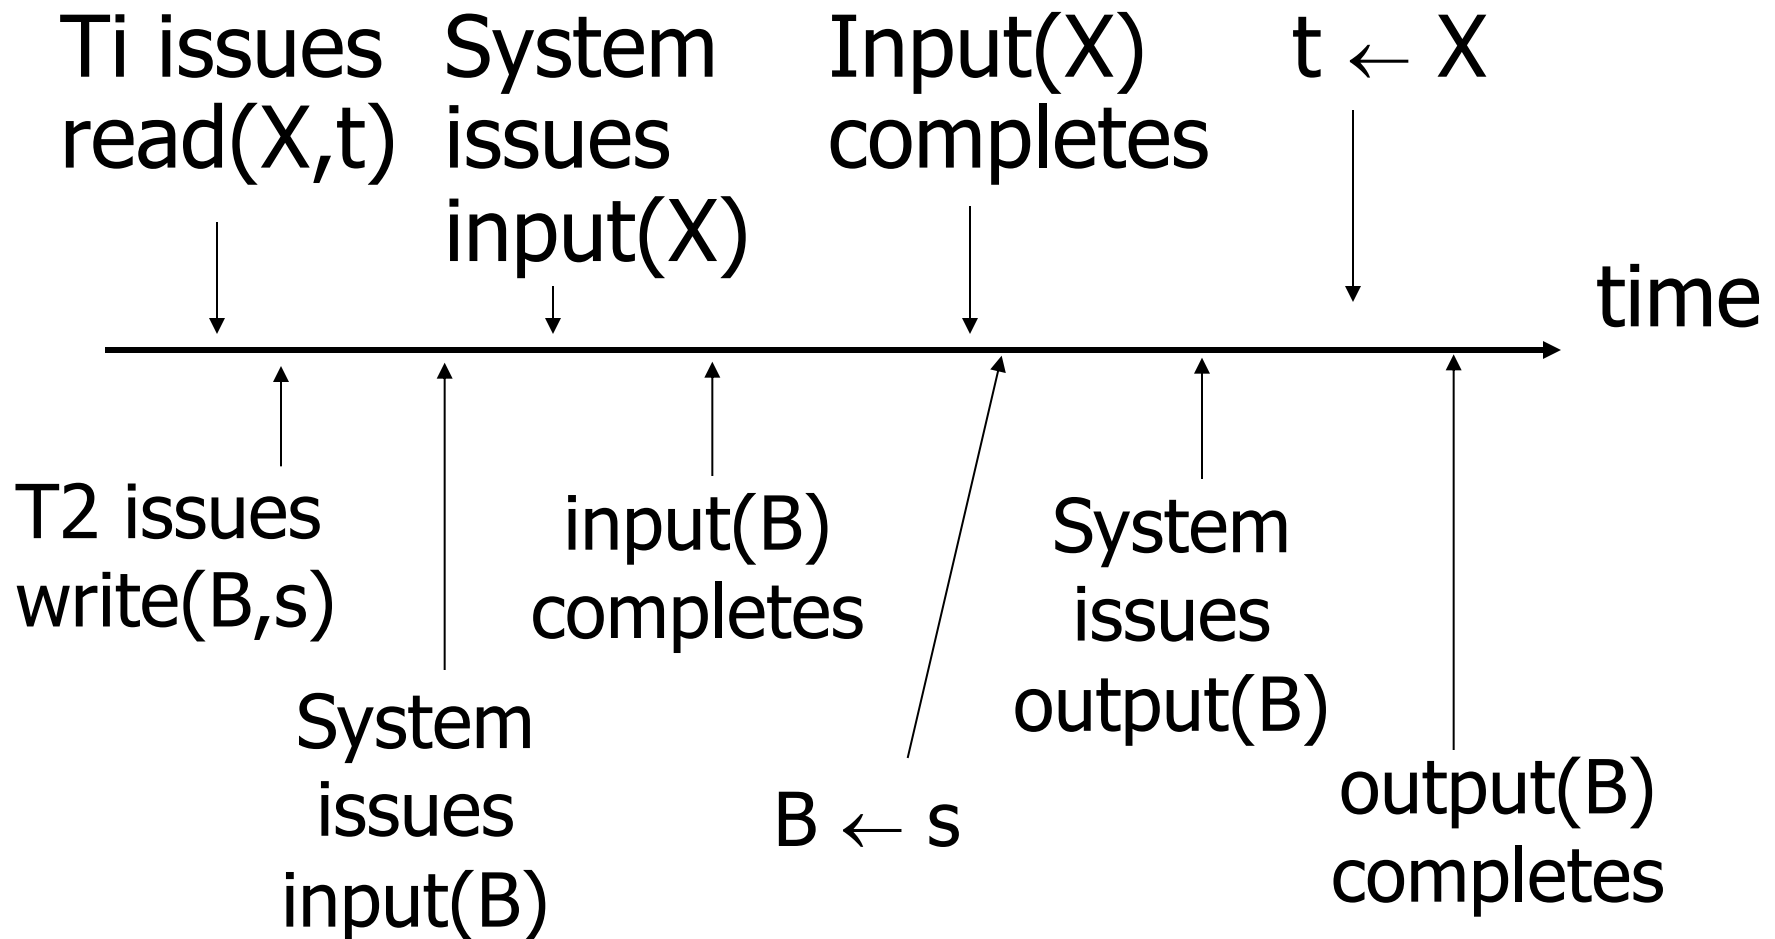


$r_1(A)$   $W_2(A)$   $W_1(A)$   
 $W_2(A)$   $r_1(A)$   $W_2(A)$

*Schedule:* represents chronological order  
in which actions are executed

*Serial schedule:* no interleaving of actions  
or transactions

# What about concurrent actions?




So net effect is either

- $S = \dots r_1(X) \dots w_2(B) \dots$  or
- $S = \dots w_2(B) \dots r_1(X) \dots$

We assume that elementary **actions are atomic**, and follow each other.

What about conflicting, concurrent actions on same object?



- Assume equivalent to either  $r_1(A) w_2(A)$   
or  $w_2(A) r_1(A)$
- $\Rightarrow$  low level synchronization mechanism
- Assumption called "atomic actions"

## Definition

$S_1, S_2$  are conflict equivalent schedules if  $S_1$  can be transformed into  $S_2$  by a series of non-conflicting swaps of adjacent actions.

A schedule is conflict serializable if it is conflict equivalent to some serial schedule.

Example:

$r_1(A); w_1(A); r_2(A); \underline{w_2(A)}; \underline{r_1(B)}; w_1(B); r_2(B); w_2(B);$

We claim this schedule is conflict-serializable.

1.  $r_1(A); w_1(A); r_2(A); \underline{w_2(A)}; \underline{r_1(B)}; w_1(B); r_2(B); w_2(B);$
2.  $r_1(A); w_1(A); \underline{r_2(A)}; \underline{r_1(B)}; w_2(A); w_1(B); r_2(B); w_2(B);$
3.  $r_1(A); w_1(A); r_1(B); r_2(A); \underline{w_2(A)}; \underline{w_1(B)}; r_2(B); w_2(B);$
4.  $r_1(A); w_1(A); r_1(B); \underline{r_2(A)}; \underline{w_1(B)}; w_2(A); r_2(B); w_2(B);$
5.  $r_1(A); w_1(A); r_1(B); w_1(B); r_2(A); w_2(A); r_2(B); w_2(B);$



- Conflict-serializability is a **sufficient** condition for serializability i.e., a conflict-serializable schedule is a serializable schedule.
- Conflict-serializability is **not required** for a schedule to be serializable.
- Schedulers **in commercial systems** generally use conflict-serializability when they need to guarantee serializability.

## Precedence graph $P(S)$ ( $S$ is schedule)

Nodes: transactions in  $S$

Arcs:  $T_i \rightarrow T_j$  whenever

- $p_i(A), q_j(A)$  are actions in  $S$
- $p_i(A) <_S q_j(A)$  ( $p_i(A)$  precedes  $q_j(A)$  in  $S$ )
- at least one of  $p_i, q_j$  is a write

## Exercise:

- What is  $P(S)$  for  
 $S = w_3(A) \ w_2(C) \ r_1(A) \ w_1(B) \ r_1(C) \ w_2(A) \ r_4(A) \ w_4(D)$
- Is  $S$  serializable?

## Another Exercise:

- What is  $P(S)$  for  
 $S = w_1(A) \ r_2(A) \ r_3(A) \ w_4(A) \ ?$

## Lemma

$S_1, S_2$  conflict equivalent  $\Rightarrow P(S_1)=P(S_2)$

# Lemma

$S_1, S_2$  conflict equivalent  $\Rightarrow P(S_1)=P(S_2)$

Proof:

Assume  $P(S_1) \neq P(S_2)$

$\Rightarrow \exists T_i: T_i \rightarrow T_j$  in  $S_1$  and not in  $S_2$

$\Rightarrow S_1 = \dots p_i(A) \dots q_j(A) \dots$	$\left\{ \begin{array}{l} p_i, q_j \\ \text{conflict} \end{array} \right.$
$S_2 = \dots q_j(A) \dots p_i(A) \dots$	

$\Rightarrow S_1, S_2$  not conflict equivalent

Note:  $P(S_1)=P(S_2) \not\Rightarrow S_1, S_2$  conflict equivalent

Counter example:

$S_1 = w_1(A) \ r_2(A) \quad w_2(B) \ r_1(B) \quad (\text{cannot swap})$

$S_2 = r_2(A) \ w_1(A) \quad r_1(B) \ w_2(B)$

# Theorem

$P(S_1)$  **acyclic**  $\iff S_1$  **conflict serializable**

( $\Leftarrow$ ) Assume  $S_1$  is conflict serializable

$\Rightarrow \exists S_s: S_s, S_1$  conflict equivalent

$\Rightarrow P(S_s) = P(S_1)$

$\Rightarrow P(S_1)$  acyclic since  $P(S_s)$  is acyclic



# Theorem

$P(S_1)$  acyclic  $\iff S_1$  conflict serializable


( $\Rightarrow$ ) Assume  $P(S_1)$  is acyclic

Transform  $S_1$  as follows:

(1) Take  $T_1$  to be transaction with no incident arcs


(2) Move all  $T_1$  actions to the front

$S_1 = \dots q_j(A) \dots p_1(A) \dots$



(3) we now have  $S_1 = \langle T_1 \text{ actions} \rangle \langle \dots \text{rest} \dots \rangle$

(4) repeat above steps to serialize rest!




# How to enforce serializable schedules?

*Option 1:* run system, recording  $P(S)$ ;  
at end of day, **check for**  $P(S)$   
**cycles** and declare if execution  
was good

# How to enforce serializable schedules?

*Option 2:* prevent P(S) cycles from occurring




# A locking protocol

Two new actions:

lock (exclusive):  $li(A)$

unlock:  $ui(A)$



## Rule #1: Consistency of transactions

$T_i: \dots l_i(A) \dots p_i(A) \dots u_i(A) \dots$

1. A transaction can only read or write an element if it previously was granted a lock on that element and hasn't yet released the lock.
2. If a transaction locks an element, it must later unlock that element.

## Rule #2

# Legality of schedules

$$S = \dots \text{li}(A) \dots \text{ui}(A) \dots$$

no  $l_j(A)$

*Locks must have their intended meaning:*  
*no two* transactions may have locked the same element without one having first released the lock.

## Exercise:

- What **schedules** are **legal**?

What **transactions** are **consistent**?

$S1 = l_1(A)l_1(B)r_1(A)w_1(B)l_2(B)u_1(A)u_1(B)$   
 $r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

$S2 = l_1(A)r_1(A)w_1(B)u_1(A)u_1(B)$   
 $l_2(B)r_2(B)w_2(B)l_3(B)r_3(B)u_3(B)$

$S3 = l_1(A)r_1(A)u_1(A)l_1(B)w_1(B)u_1(B)$   
 $l_2(B)r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

## Exercise:

- What schedules are legal?

What transactions are consistent?

$S1 = l_1(A)l_1(B)r_1(A)w_1(B)l_2(B)u_1(A)u_1(B)$   
 $r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$  (S1: not legal)

$S2 = l_1(A)r_1(A)w_1(B)u_1(A)u_1(B)$  (T1: not consistent)

$l_2(B)r_2(B)w_2(B)l_3(B)r_3(B)u_3(B)$  (S2: not legal)

$S3 = l_1(A)r_1(A)u_1(A)l_1(B)w_1(B)u_1(B)$

$l_2(B)r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

(S3: legal schedule, T1,T2,T3: consistent transactions)




# Schedule F (legal schedule of consistent transactions) (still not equivalent to a serial schedule)


		A	B
T1	T2	25	25
l <sub>1</sub> (A);Read(A)			
A ← A+100;Write(A);u <sub>1</sub> (A)		125	
	l <sub>2</sub> (A);Read(A)		
	A ← Ax2;Write(A);u <sub>2</sub> (A)	250	
	l <sub>2</sub> (B);Read(B)		
	B ← Bx2;Write(B);u <sub>2</sub> (B)		50
l <sub>1</sub> (B);Read(B)			
B ← B+100;Write(B);u <sub>1</sub> (B)			150
		250	150

# Rule #3 Two phase locking (2PL) for transactions


$T_i = \dots \dots \text{li}(A) \dots \dots \text{ui}(A) \dots \dots$



no unlocks



no locks



# Schedule G


T1	T2
$l_1(A); \text{Read}(A)$	
$A \leftarrow A + 100; \text{Write}(A)$	
$l_1(B); u_1(A)$	
	$l_2(A); \text{Read}(A)$
	$A \leftarrow A \times 2; \text{Write}(A); l_2(B)$

delayed

# Schedule G

T1	T2
$l_1(A); \text{Read}(A)$	
$A \leftarrow A + 100; \text{Write}(A)$	
$l_1(B); u_1(A)$	
	$l_2(A); \text{Read}(A)$
	$A \leftarrow A \times 2; \text{Write}(A); l_2(B)$
$\text{Read}(B); B \leftarrow B + 100$	
$\text{Write}(B); u_1(B)$	

delayed



# Schedule G

T1	T2
$l_1(A); \text{Read}(A)$	
$A \leftarrow A + 100; \text{Write}(A)$	
$l_1(B); u_1(A)$	
	$l_2(A); \text{Read}(A)$ delayed
	$A \leftarrow A \times 2; \text{Write}(A); l_2(B)$
$\text{Read}(B); B \leftarrow B + 100$	
$\text{Write}(B); u_1(B)$	
	$l_2(B); u_2(A); \text{Read}(B)$
	$B \leftarrow B \times 2; \text{Write}(B); u_2(B);$

# Schedule G "good" (equivalent to a serial)

		A	B	
T1	T2	25	25	
l <sub>1</sub> (A);Read(A); A←A+100		125		
Write(A); l <sub>1</sub> (B);u <sub>1</sub> (A)				
	l <sub>2</sub> (A);Read(A)	250		
	A←Ax2;Write(A);u <sub>2</sub> (A)			
	l <sub>2</sub> (B); delayed		125	
Read(B);				
B←B+100;Write(B);u <sub>1</sub> (B)				
	l <sub>2</sub> (B);Read(B)	250	250	
	B←Bx2;Write(B);u <sub>2</sub> (B)			

Theorem Rules #1,2,3  $\Rightarrow$  conflict  
(consistency, legality, 2PL) serializable  
schedule



Theorem Rules #1,2,3  $\Rightarrow$  conflict  
(consistency, legality, 2PL) serializable  
schedule

Intuitively, each two-phase-locked transaction may be thought to execute in its entirety at the instant it issues its first unlock request. In a conflict-equivalent serial schedule transactions appear in the same order as their first unlocks.

Theorem Rules #1,2,3  $\Rightarrow$  conflict  
serializable  
schedule

Proof:

**BASIS:** If  $n = 1$ , there is nothing to do;  $S$  is already a serial schedule.

**INDUCTION:** Suppose  $S$  involves  $n$  transactions  $T_1, T_2, \dots, T_n$ , and let  $T_i$  be the transaction with the first unlock action in the entire schedule  $S$ , say  $u_i(x)$ .

We claim it is possible to move all the read and write actions of  $T_i$  forward to the beginning of the schedule without passing any conflicting reads or writes.

Theorem Rules #1,2,3  $\Rightarrow$  conflict  
(2PL) serializable  
schedule

To help in proof:

Definition Shrink( $T_i$ ) = SH( $T_i$ ) =  
first unlock action of  $T_i$

## Lemma


$$T_i \rightarrow T_j \text{ in } S \Rightarrow SH(T_i) <_S SH(T_j)$$

Proof of lemma:

$T_i \rightarrow T_j$  means that

$$S = \dots p_i(A) \dots q_j(A) \dots; \quad p, q \text{ conflict}$$

By rules 1,2:

$$S = \dots p_i(A) \dots u_i(A) \dots l_j(A) \dots q_j(A) \dots$$


By rule 3:

$$SH(T_i) \qquad SH(T_j)$$

$$\text{So, } SH(T_i) <_S SH(T_j)$$

Theorem Rules #1,2,3  $\Rightarrow$  conflict  
(2PL) serializable  
schedule

Proof:

(1) Assume  $P(S)$  has cycle


$$T_1 \rightarrow T_2 \rightarrow \dots T_n \rightarrow T_1$$

(2) By lemma:  $SH(T_1) < SH(T_2) < \dots < SH(T_1)$

(3) Impossible, so  $P(S)$  acyclic


(4)  $\Rightarrow S$  is conflict serializable

# Schedule H (T<sub>2</sub> reversed)



Deadlock !!!

- Assume **deadlocked transactions** are **rolled back**
  - They have no effect
  - They do not appear in schedule

E.g., Schedule H =   
This space intentionally  
left blank!


# Deadlocks

- Detection
  - Wait-for graph
- Prevention
  - Resource ordering
  - Timeout
  - Wait-die
  - Wound-wait



# Deadlock Detection

- Build **Wait-For graph** ( $T_i \rightarrow T_j$  edge if  $T_i$  waits for  $T_j$ )
- Use lock table structures
- Build incrementally or periodically
- When cycle found, **rollback victim**



## Resource Ordering (prevention)

- Order all elements  $A_1, A_2, \dots, A_n$
- A transaction  $T$  can lock  $A_i$  after  $A_j$  only if  $i > j$


Problem : Ordered lock requests **not realistic** in most cases

# Timeout

- If transaction **waits** more than **L sec.**,  
roll it back!
- Simple scheme
- Hard to select L

# 2PL subset of Serializable

(schedules that can be implemented by 2PL locks are subset of serializable schedules)



S1: w1(x) w3(x) w2(y) w1(y)

- **S1 cannot be achieved via 2PL:**  
The lock by T1 for y must occur after w2(y), so the unlock by T1 for x must occur after this point (and before w3(x)). Thus, w3(x) cannot occur under 2PL where shown in S1 because T1 holds the x lock at that point.
- However, **S1 is serializable** (conflict-serializable) (equivalent to T2, T1, T3).

If you need a bit more practice:  
Are our schedules  $S_C$  and  $S_D$  2PL schedules?

$S_C$ :  $w1(A)$   $w2(A)$   $w1(B)$   $w2(B)$

$S_D$ :  $w1(A)$   $w2(A)$   $w2(B)$   $w1(B)$

- Beyond this simple 2PL protocol, it is all a matter of improving performance and allowing more concurrency....
  - Shared locks
  - Multiple granularity
  - Inserts, deletes and phantoms
  - Other types of C.C. mechanisms

# Shared locks

So far:

$S = \dots l_1(A) \ r_1(A) \ u_1(A) \ \dots \ l_2(A) \ r_2(A) \ u_2(A) \ \dots$

Do not conflict





## Shared locks

So far:

$S = \dots l_1(A) \ r_1(A) \ u_1(A) \ \dots \ l_2(A) \ r_2(A) \ u_2(A) \ \dots$

Do not conflict



Instead:

$S = \dots \textcolor{red}{l}_{s1}(A) \ r_1(A) \ \textcolor{red}{l}_{s2}(A) \ r_2(A) \ \dots \ u_{s1}(A) \ u_{s2}(A)$

## Lock actions

$l-t_i(A)$ : lock A in t mode (**t is S or X**)

$u-t_i(A)$ : unlock t mode (t is S or X)

## Shorthand:

$u_i(A)$ : unlock whatever modes

$T_i$  has locked A

## Rule #1 Consistency of transactions

$T_i = \dots I-S_1(A) \dots r_1(A) \dots u_1(A) \dots$

$T_i = \dots I-X_1(A) \dots w_1(A) \dots u_1(A) \dots$

- A transaction **may not write without holding an exclusive lock**, and **may not read without holding some lock**.
- All locks must be followed by an unlock of the same element.

## Two-phase locking of transactions:

Locking must precede unlocking.

## Legality of schedules:

An element may either be locked exclusively by one transaction or by several in shared mode, but not both.

- What about transactions that **read and write same object**?

Option 1: Request **exclusive** lock

$T_i = \dots \text{I-X}_1(A) \dots r_1(A) \dots w_1(A) \dots u(A) \dots$

- What about transactions that read and write same object?

## Option 2: Upgrade

(E.g., need to read, but don't know if will write...)

$T_i = \dots I-S_1(A) \dots r_1(A) \dots I-X_1(A) \dots w_1(A) \dots u(A) \dots$



Think of

- Get 2nd lock on A, or
- Drop S, get X lock


## Rule #2 Legal scheduler

$$S = \dots l-S_i(A) \dots \dots u_i(A) \dots$$



no  $l-X_j(A)$

$$S = \dots l-X_i(A) \dots \dots u_i(A) \dots$$



no  $l-X_j(A)$

no  $l-S_j(A)$

## A way to summarize Rule #2

### Compatibility matrix

Comp

	S	X
S	true	false
X	false	false



## Rule # 3      2PL transactions

No change except for **upgrades**:

(I) If upgrade gets more locks

(e.g.,  $S \rightarrow \{S, X\}$ ) then no change!

(It is like a new lock, allowed only in the growing phase)

(II) If upgrade releases read (shared) lock (e.g.,  $S \rightarrow X$ )

- can be allowed in growing phase

(This is not a real release or unlock)

Theorem Rules 1,2,3  $\Rightarrow$  Conf.serializable  
for S/X locks schedules

Proof: similar to X locks case

Detail:

$l-t_i(A), l-r_j(A)$  do not conflict if  $\text{comp}(t,r)$

$l-t_i(A), u-r_j(A)$  do not conflict if  $\text{comp}(t,r)$


## Lock types beyond S/X

Examples:

- (1) increment lock
- (2) update lock

## Example (1): increment lock

- Atomic increment action:  $IN_i(A)$   
 $\{Read(A); A \leftarrow A+k; Write(A)\}$
- $IN_i(A), IN_j(A)$  do not conflict!



Comp


	S	X	I
S			
X			
I			

Comp

	S	X	I
S	T	F	F
X	F	F	F
I	F	F	T

# Update locks

A common **deadlock** problem **with upgrades**:



## Solution

If  $T_i$  wants to read  $A$  and knows it may later want to write  $A$ , it requests update lock (not shared)



## New request

Comp

Lock  
already  
held in

	S	X	U
S	T	F	T
X	F	F	F
U	F	F	F

-> not symmetric table

Note: object A may be locked in different modes at the same time...

$$S_1 = \dots I-S_1(A) \dots I-S_2(A) \dots I-U_3(A) \dots \left\{ \begin{array}{l} I-S_4(A) \dots ? \\ I-U_4(A) \dots ? \end{array} \right.$$


- To grant a lock in mode t, mode t **must be compatible with all currently held locks** on object

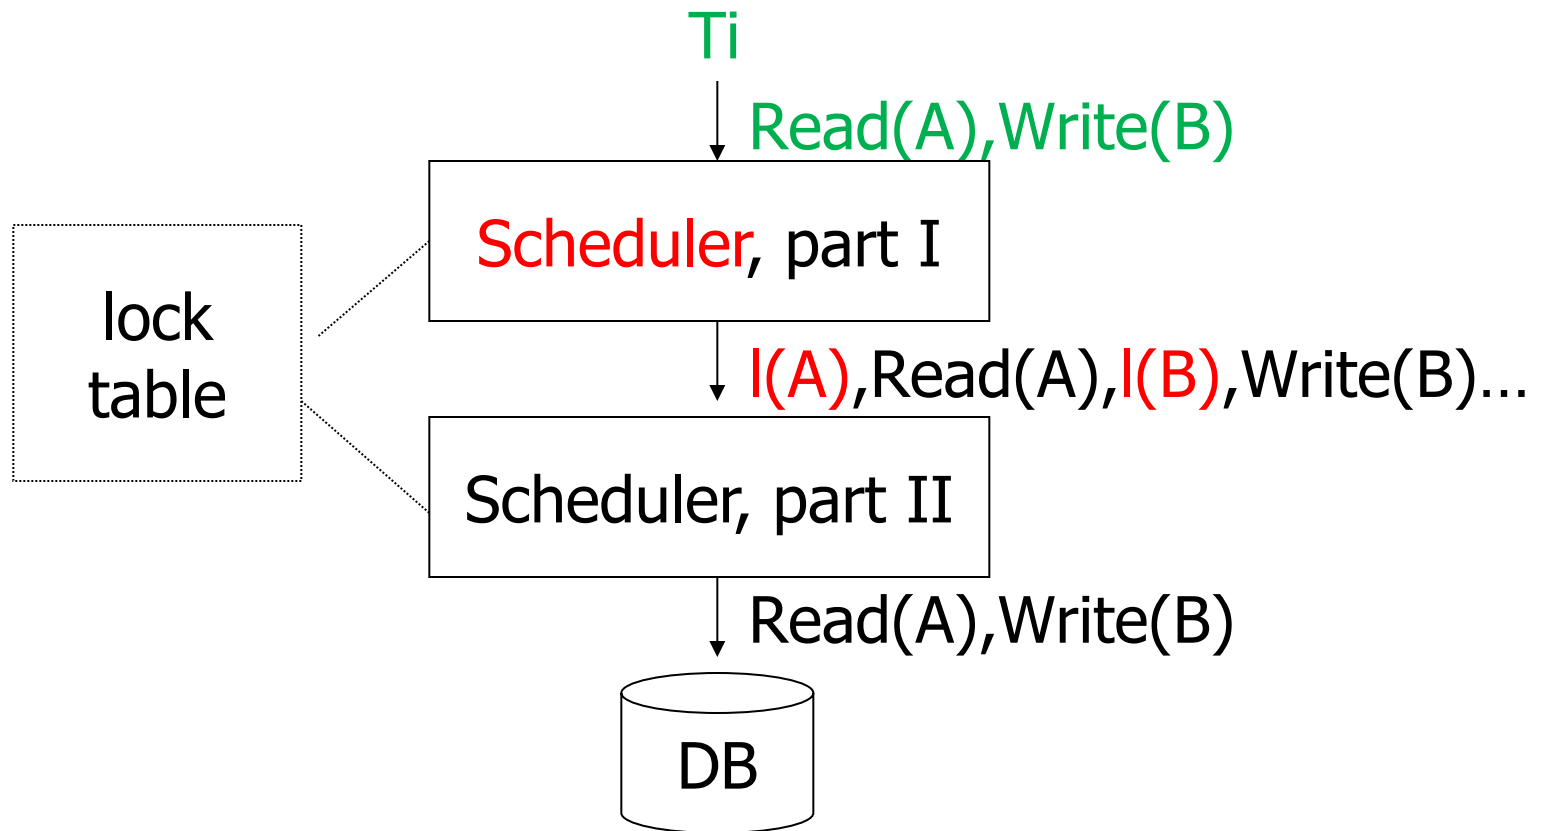
# How does **locking** work **in practice**?

- Every system is different  
(E.g., may not even provide  
CONFLICT-SERIALIZABLE schedules)
- But here is one (simplified) way ...

## Sample Locking System:

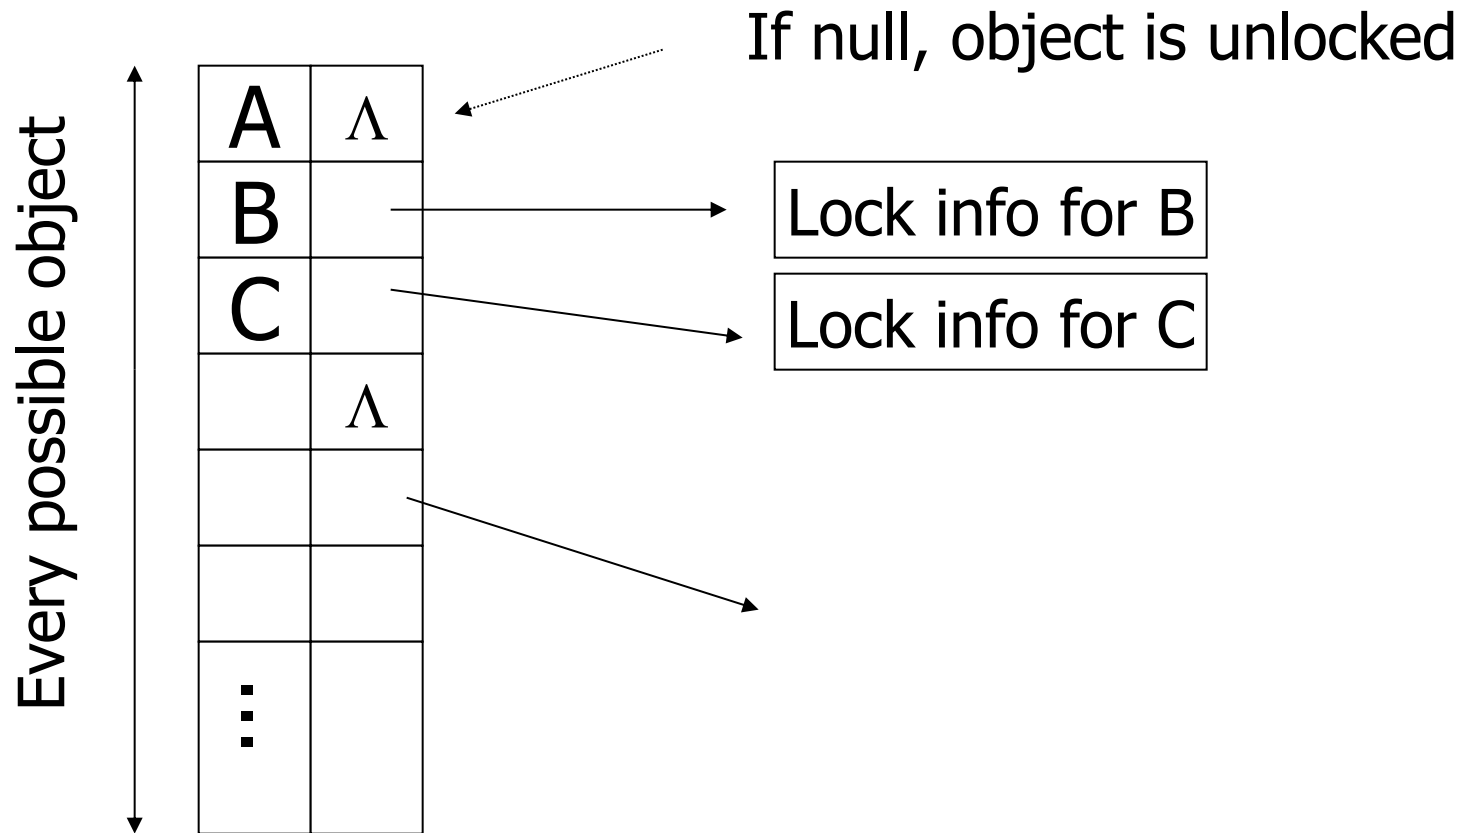
- (1) Don't trust transactions to request/release locks
- (2) Hold all locks until transaction commits






Scheduler requests locks (not transactions)

# Lock table      Conceptually

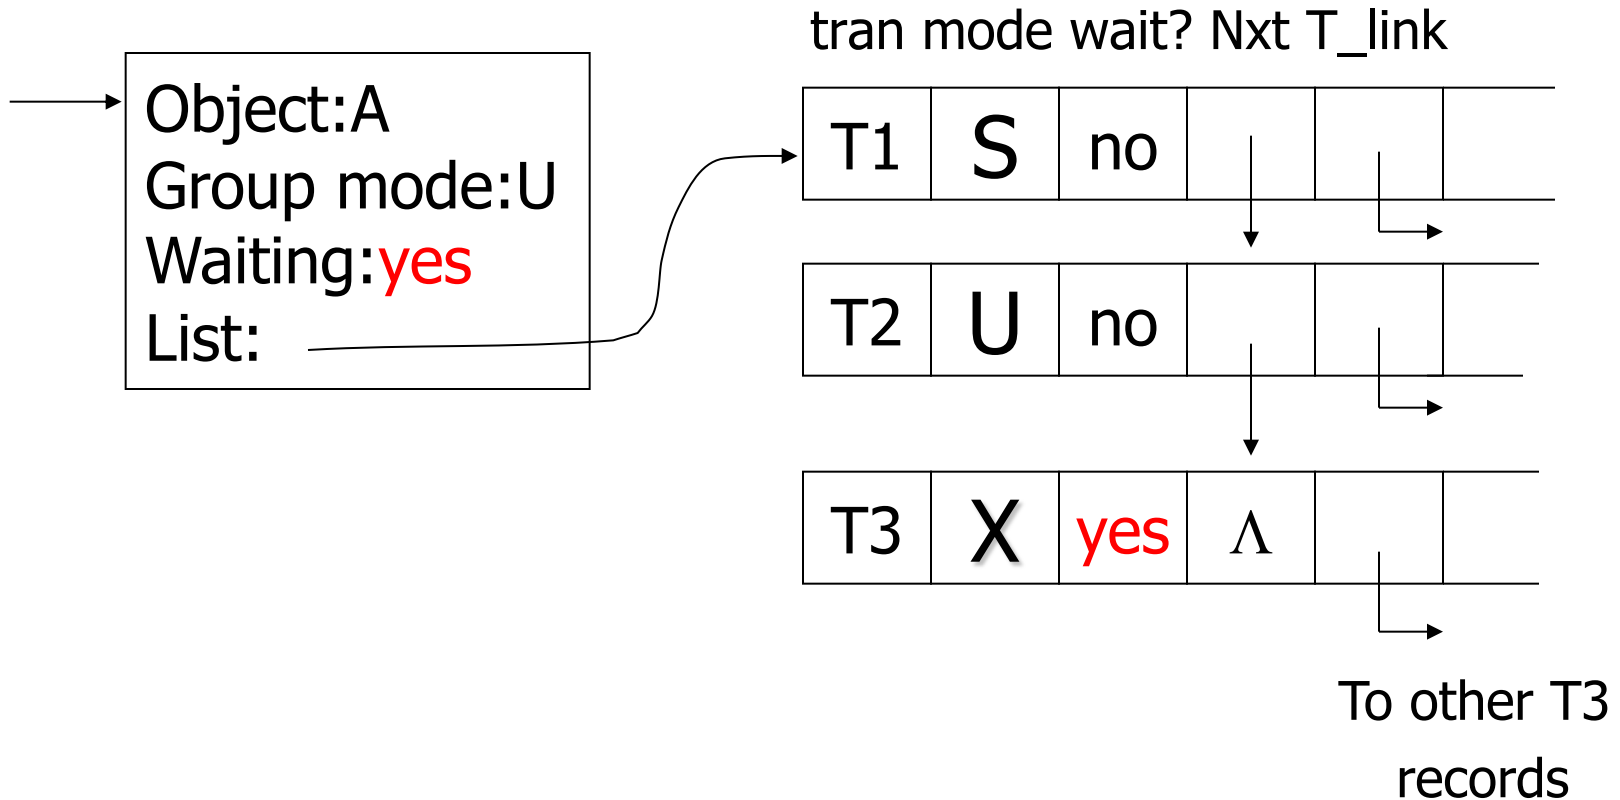


But use hash table:




If object not found in hash table, it is unlocked

# Lock info for A - example







# What are the objects we lock?



DB



DB



DB

?


- Locking works in any case, but should we choose small or large objects?

- Locking works in any case, but should we choose small or large objects?
- If we lock large objects (e.g., Relations)
  - Need few locks
  - Low concurrency
- If we lock small objects (e.g., tuples, fields)
  - Need more locks
  - More concurrency


# We can have it both ways!!

## Ask any janitor to give you the solution...


Should we close (lock) individual doors or restroom?



# Example (R: relation, t: tuple)




## Example (Warning or **Intention**)




**IS:** Intention to obtain a shared lock on a subelement

# Example




## Example (b)





## Example (b)



# Multiple granularity (SIX: group mode)

Comp

Requestor

		IS	IX	S	SIX	X
Holder	IS					
	IX					
	S					
	SIX					
	X					


# Multiple granularity (SIX: group mode)

Comp


Requestor

		IS	IX	S	SIX	X
Holder	IS	T	T	T	T	F
	IX	T	T	F	F	F
	S	T	F	T	F	F
	SIX	T	F	F	F	F
	X	F	F	F	F	F

Parent locked in	Child can be locked in
IS	
IX	
S	
SIX	
X	



Parent locked in	Child can be locked by same transaction in
IS	IS, S
IX	IS, S, IX, X, SIX
S	none
SIX	X, IX, [SIX]
X	none



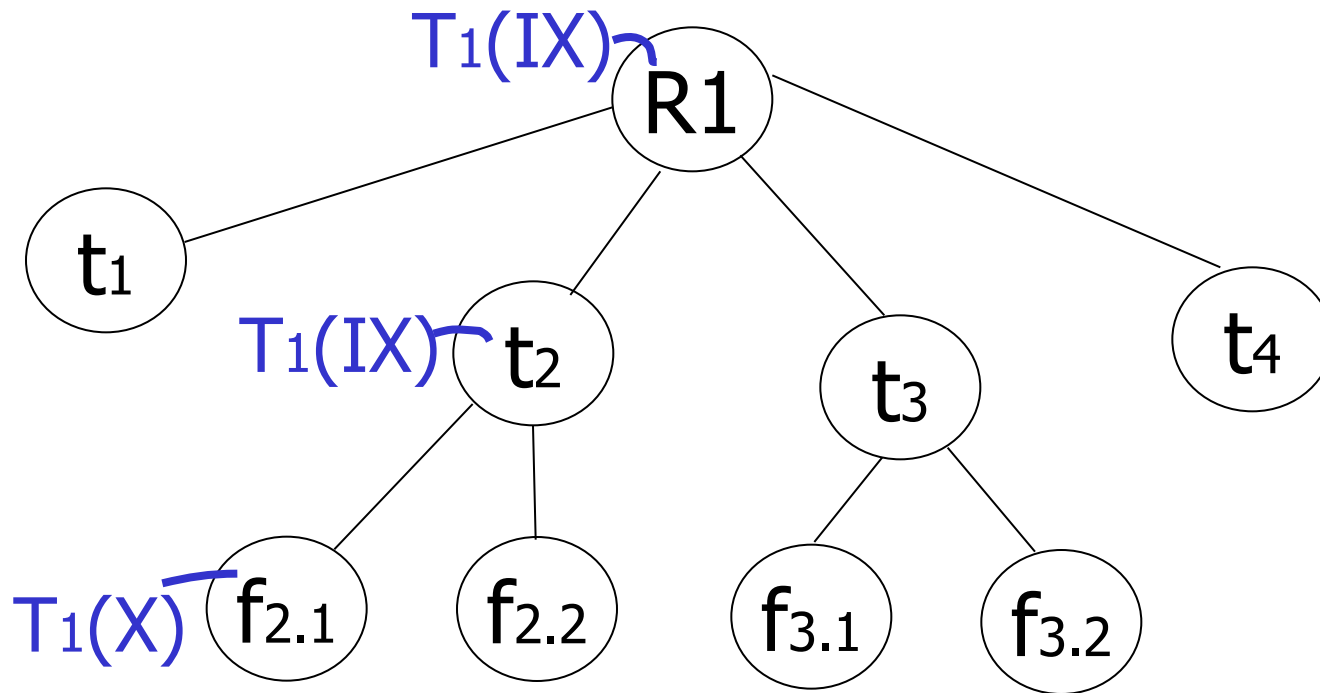
not necessary

# Rules

- (1) Follow multiple granularity compat. function
- (2) **Lock root** of tree **first**, any mode
- (3) Node Q can be locked by Ti in S or IS only if parent(Q) locked by Ti in IX or IS
- (4) Node Q can be locked by Ti in X,SIX,IX only if parent(Q) locked by Ti in IX,SIX
- (5) Ti is two-phase
- (6) Ti can **unlock** node Q **only if none of Q's children are locked** by Ti


## Exercise:

- Can T2 access object f2.2 in X mode?  
What locks will T2 get?



## Exercise:


- Can T2 access object f2.2 in X mode?  
What locks will T2 get?






## Exercise:

- Can T2 access object f3.1 in X mode?  
What locks will T2 get?



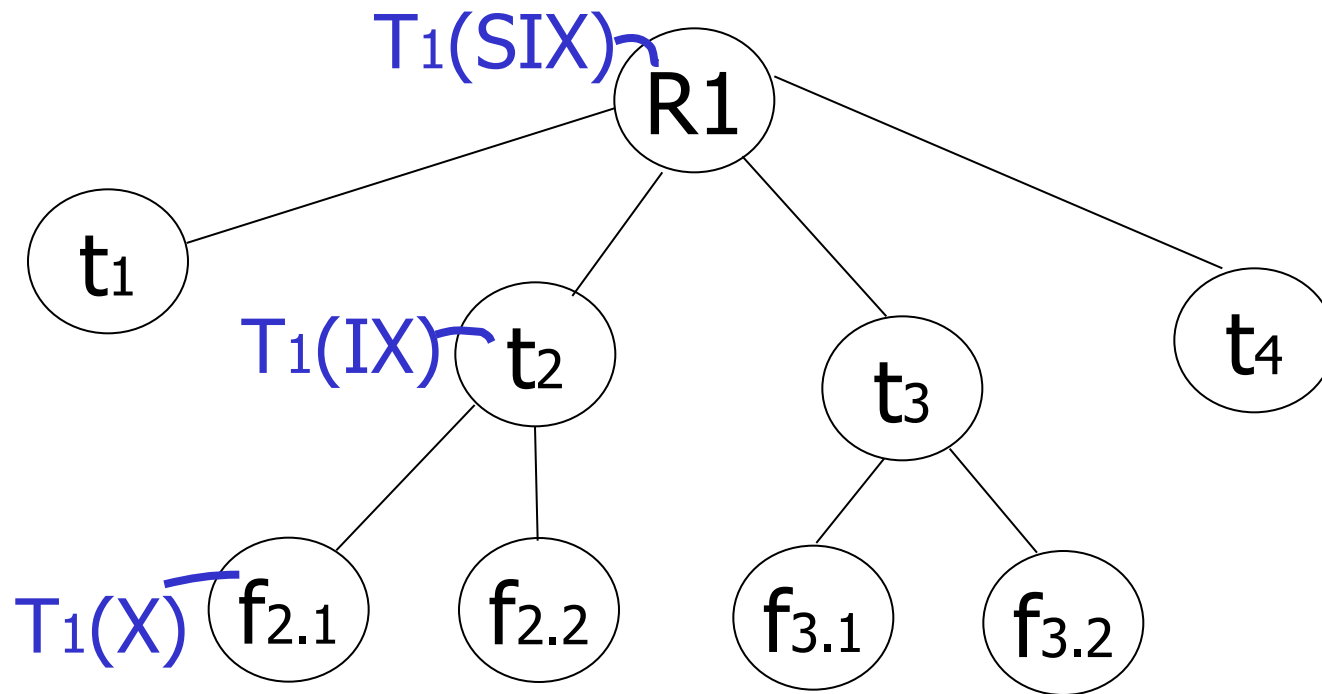
## Exercise:

- Can T2 access object f2.2 in S mode?  
What locks will T2 get?



## Exercise:

- Can T2 access object f2.2 in X mode?  
What locks will T2 get?



# Insert + delete operations

A
$\vdots$
Z
$\alpha$

← Insert

## Modifications to locking rules:

- (1) Get exclusive lock on A before deleting A
- (2) At insert A operation by  $T_i$ ,  
 $T_i$  is given exclusive lock on A

# Still have a problem: **Phantoms**

Example: relation R (E#,name,...)

constraint: E# is key

use tuple locking

R	E#	Name	....
o1	55	Smith	
o2	75	Jones	


$T_1$ : Insert  $\langle 08, \text{Obama}, \dots \rangle$  into R

$T_2$ : Insert  $\langle 08, \text{McCain}, \dots \rangle$  into R

$T_1$	$T_2$
$S_1(o_1)$	$S_2(o_1)$
$S_1(o_2)$	$S_2(o_2)$
Check Constraint	Check Constraint
$\vdots$	$\vdots$
Insert $o_3[08, \text{Obama}, \dots]$	Insert $o_4[08, \text{McCain}, \dots]$

# Solution

- Use multiple granularity tree
- Before insert of node Q,  
lock parent(Q) in  
X mode





# Back to example

T<sub>1</sub>: Insert<08,Obama>

T<sub>1</sub>

X<sub>1</sub>(R)

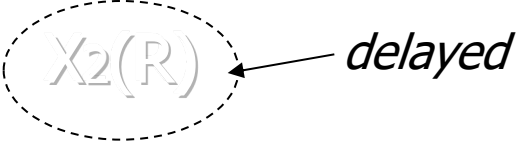
Check constraint

Insert<08,Obama>

U(R)

T<sub>2</sub>: Insert<08,McCain>

T<sub>2</sub>

X<sub>2</sub>(R) ← *delayed*

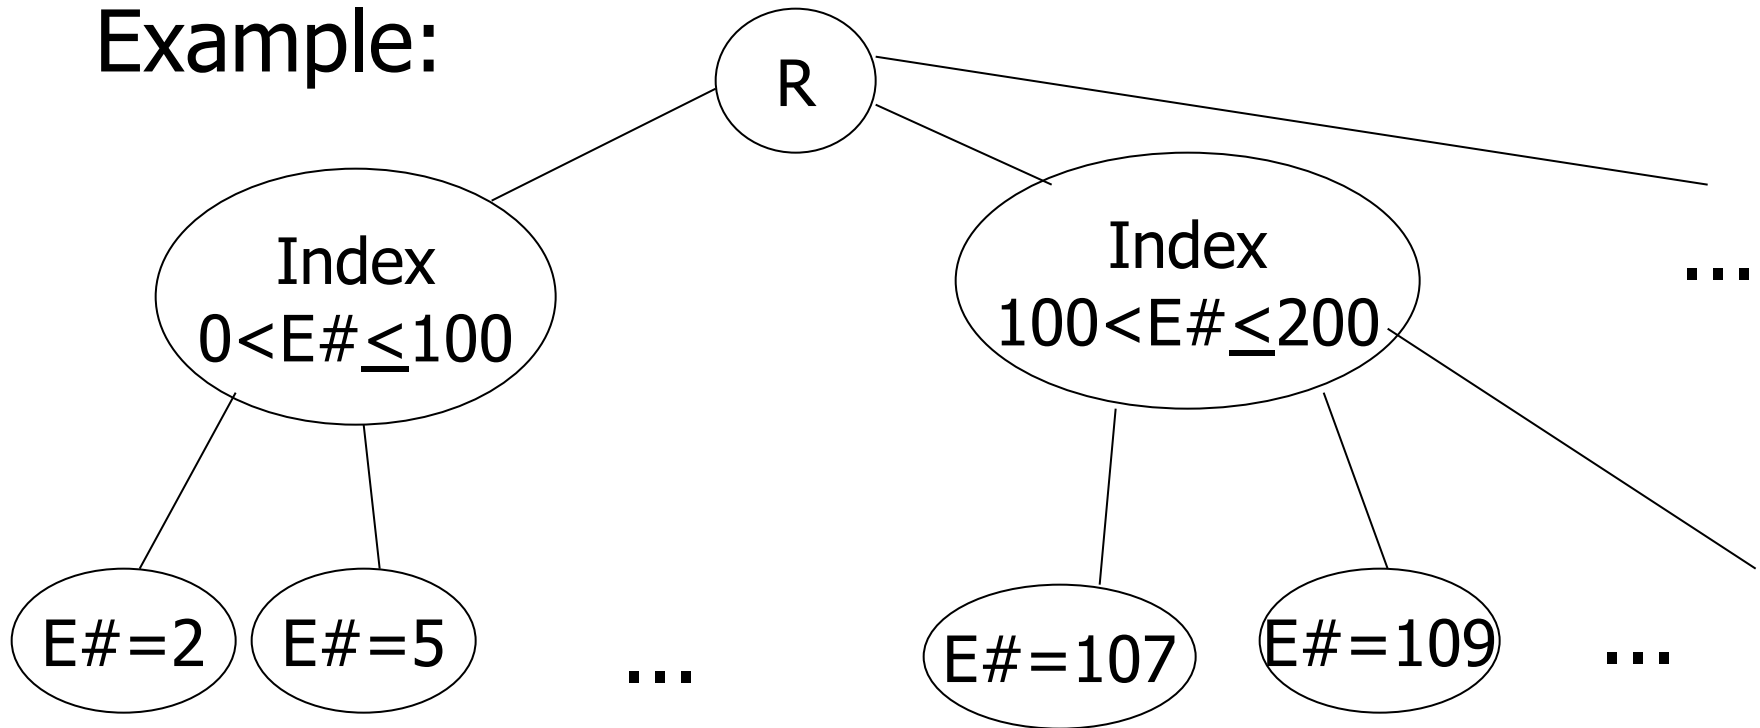
X<sub>2</sub>(R)

Check constraint

Oops! e# = 08 already in R!

# Instead of using R, can use index on R:

Example:




- This approach can be generalized to multiple indexes...

# Next:

- Tree-based concurrency control
- Validation concurrency control


# Example

- all objects accessed through root, following pointers




# Example

- all objects accessed through root, following pointers





# Example

- all objects accessed through root, following pointers





- can we release A lock if we no longer need A??

# Idea: traverse like “Monkey Bars”







# Idea: traverse like “Monkey Bars”




# Idea: traverse like “Monkey Bars”



# Why does this work?

- Assume all  $T_i$  start at root; exclusive lock
- $T_i \rightarrow T_j \Rightarrow T_i$  locks root before  $T_j$




- Actually works if we don't always start at root

## Rules: tree protocol (exclusive locks)

- (1) First lock by  $T_i$  may be on any item
- (2) After that, item  $Q$  can be locked by  $T_i$   
only if  $\text{parent}(Q)$  locked by  $T_i$
- (3) Items may be unlocked at any time
- (4) After  $T_i$  unlocks  $Q$ , it cannot relock  $Q$


- Tree-like protocols are used typically for B-tree concurrency control



E.g., during insert, do not release parent lock, until you are certain child does not have to split

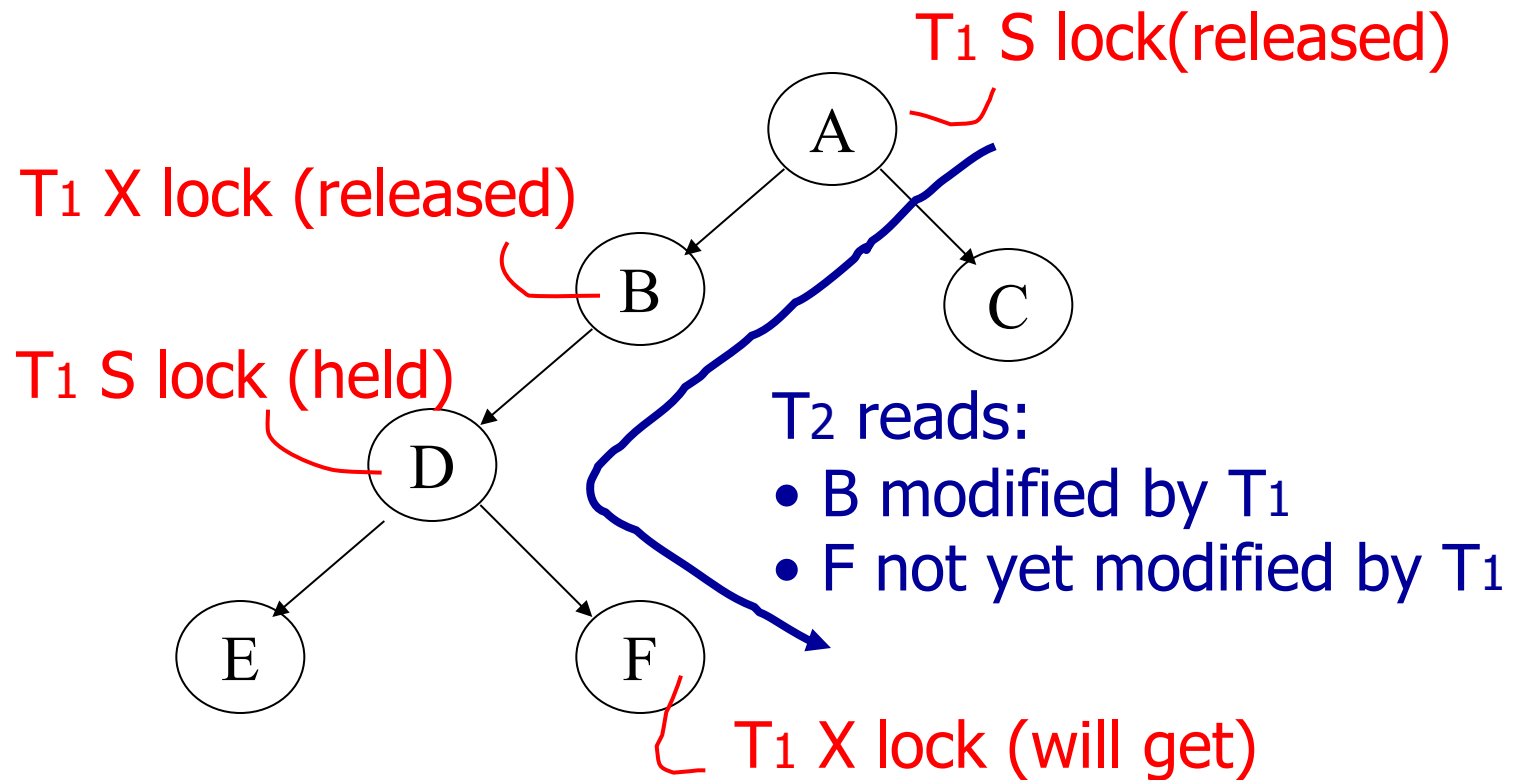
# Tree Protocol with Shared Locks

- Rules for shared & exclusive locks?



# Tree Protocol with Shared Locks

- Rules for shared & exclusive locks?



# Tree Protocol with Shared Locks

- Need more restrictive protocol
- Will this work??
  - Once  $T_1$  locks one object in X mode, all further locks down the tree must be in X mode



# Validation

Transactions have 3 phases:

## (1) Read

- all DB values read
- writes to temporary storage
- no locking

## (2) Validate

- check if schedule so far is serializable

## (3) Write

- if validate ok, write to DB

## Key idea

- Make validation atomic
- If  $T_1, T_2, T_3, \dots$  is validation order, then resulting schedule will be conflict equivalent to  $S_s = T_1 T_2 T_3 \dots$

To implement validation, system keeps two sets:

- FIN = transactions that have finished phase 3 (and are all done)
- VAL = transactions that have successfully finished phase 2 (validation)


## Example of what validation must prevent:

$$RS(T_2) = \{B\}$$

$$WS(T_2) = \{B, D\}$$

$$RS(T_3) = \{A, B\} \neq \phi$$

$$WS(T_3) = \{C\}$$



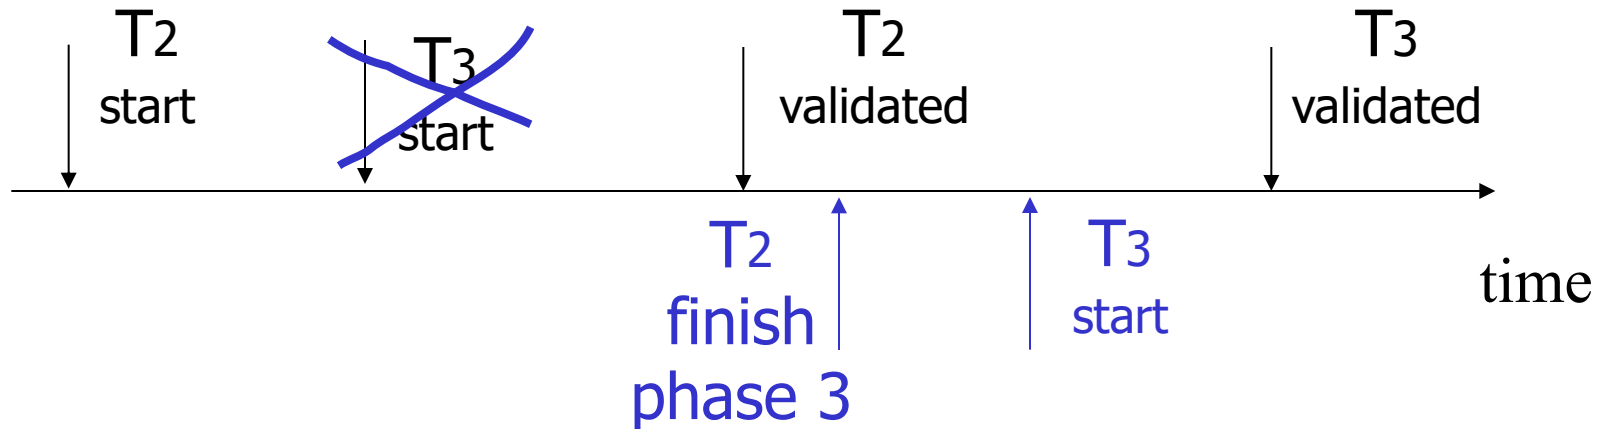
Example of what validation must ~~allow~~ prevent:

$RS(T_2) = \{B\}$

$WS(T_2) = \{B, D\}$

$RS(T_3) = \{A, B\} \neq \phi$

$WS(T_3) = \{C\}$




## Another thing validation must prevent:

$$RS(T_2) = \{A\}$$

$$RS(T_3) = \{A, B\}$$

$$WS(T_2) = \{D, E\}$$

$$WS(T_3) = \{C, D\}$$



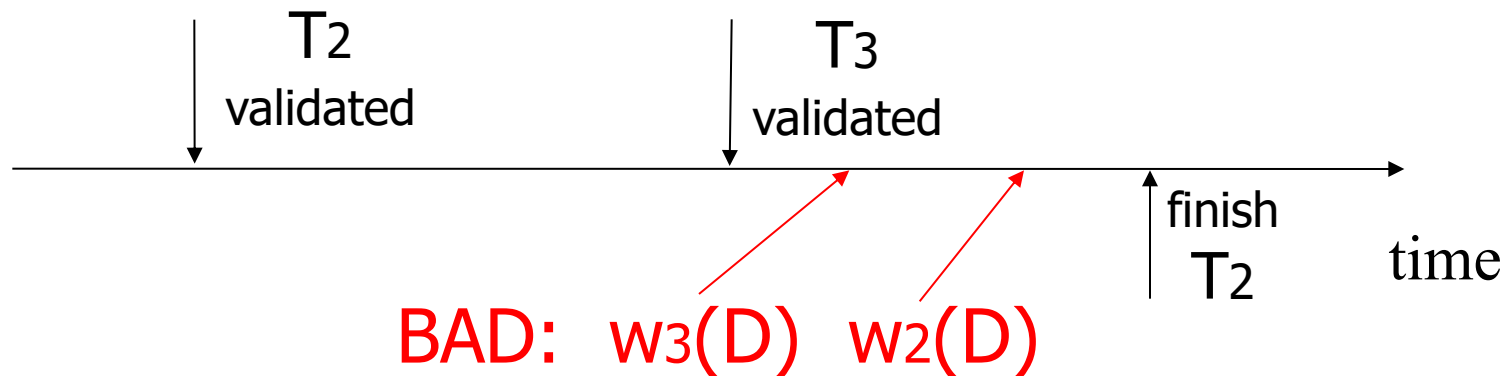
## Another thing validation must prevent:

$$RS(T_2) = \{A\}$$

$$RS(T_3) = \{A, B\}$$

$$WS(T_2) = \{D, E\}$$

$$WS(T_3) = \{C, D\}$$



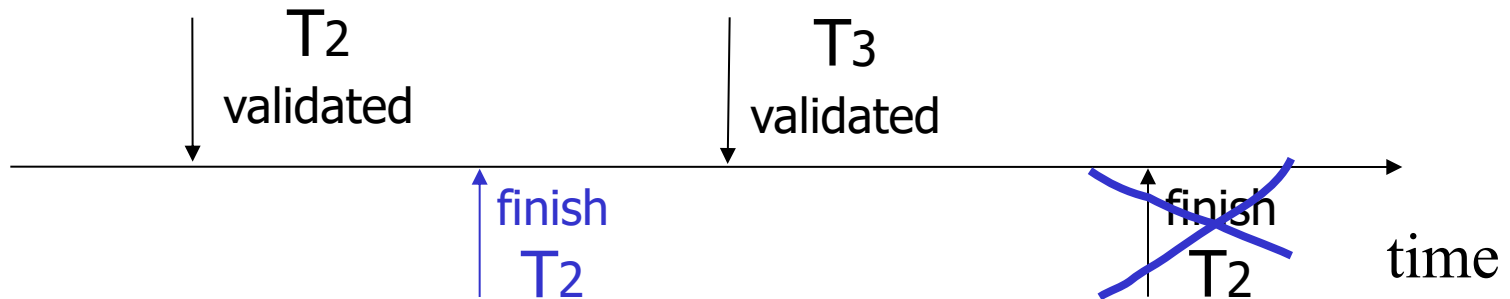
Another thing validation must ~~prevent~~ <sup>allow</sup>:

$RS(T_2) = \{A\}$

$RS(T_3) = \{A, B\}$

$WS(T_2) = \{D, E\}$

$WS(T_3) = \{C, D\}$





## Validation rules for $T_j$ :

(1) When  $T_j$  starts phase 1:

$\text{ignore}(T_j) \leftarrow \text{FIN}$

(2) at  $T_j$  Validation:

if check ( $T_j$ ) then

$[ \text{VAL} \leftarrow \text{VAL} \cup \{T_j\};$

do write phase;

$\text{FIN} \leftarrow \text{FIN} \cup \{T_j\} ]$

Check ( $T_j$ ):

For  $T_i \in \text{VAL} - \text{IGNORE}(T_j)$  DO

IF [  $\text{WS}(T_i) \cap \text{RS}(T_j) \neq \emptyset$  OR

$T_i \notin \text{FIN}$  ] THEN RETURN false;

RETURN true;

Check ( $T_j$ ):

For  $T_i \in \text{VAL} - \text{IGNORE}(T_j)$  DO

IF [  $\text{WS}(T_i) \cap \text{RS}(T_j) \neq \emptyset$  OR

$T_i \notin \text{FIN}$  ] THEN RETURN false;

RETURN true;

Is this check too restrictive ?

## Improving Check( $T_j$ )

For  $T_i \in \text{VAL} - \text{IGNORE}(T_j)$  DO

IF [  $\text{WS}(T_i) \cap \text{RS}(T_j) \neq \emptyset$  OR

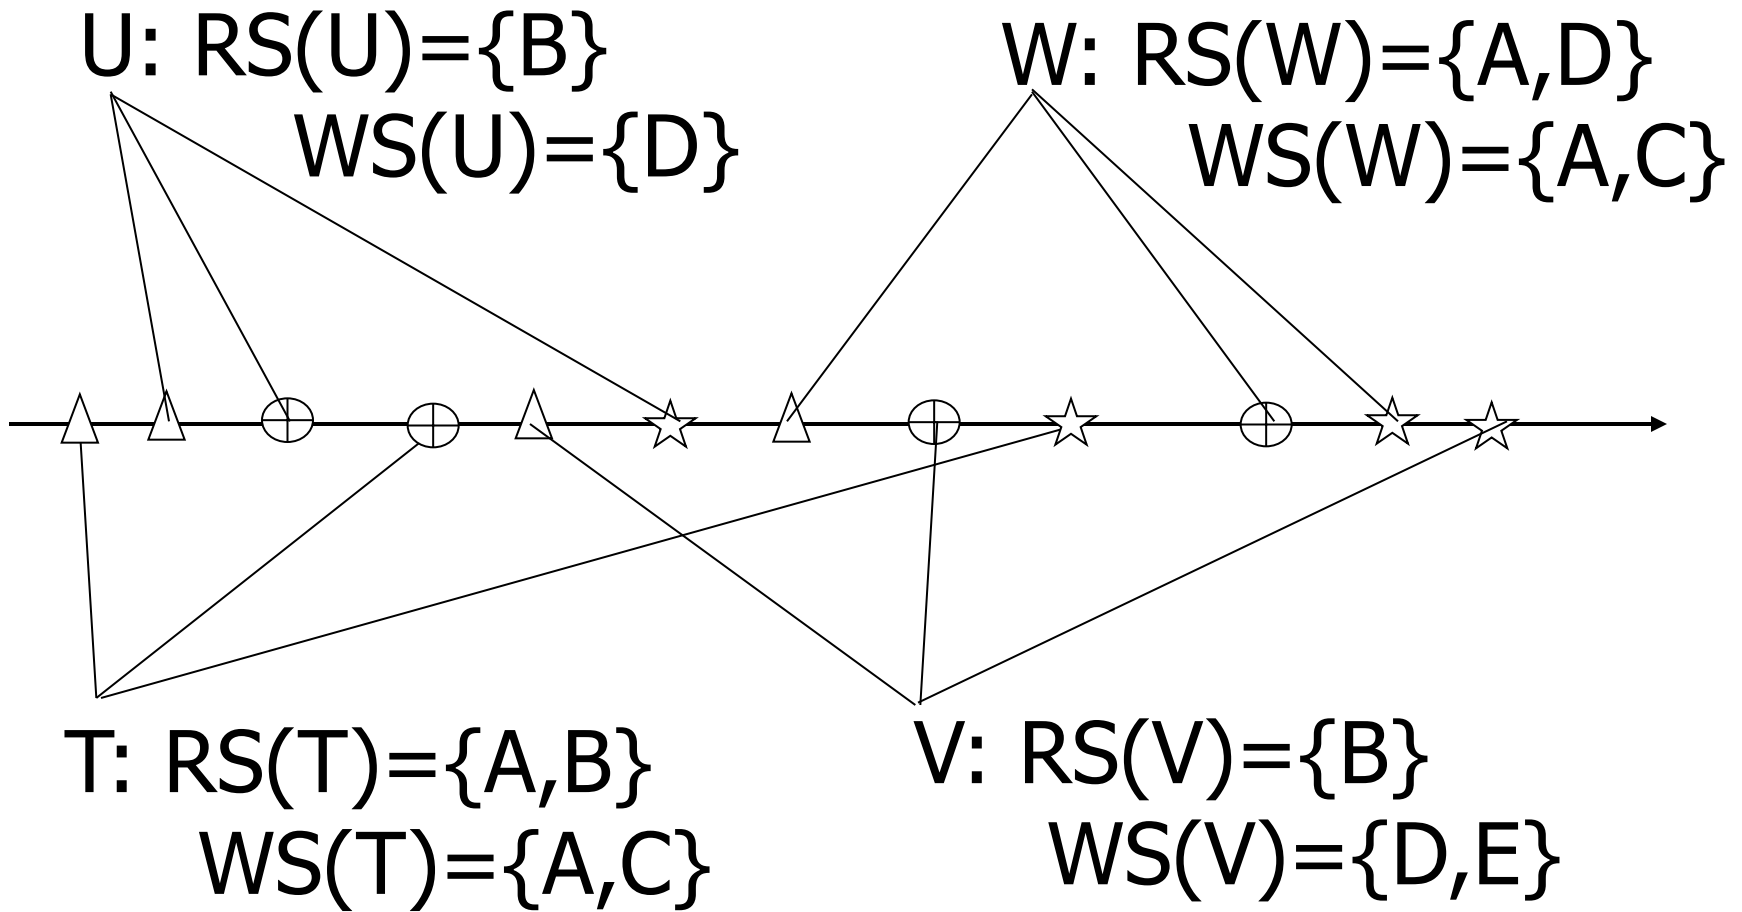
$(T_i \notin \text{FIN} \text{ AND } \text{WS}(T_i) \cap \text{WS}(T_j) \neq \emptyset)$  ]

THEN RETURN false;


RETURN true;

# Exercise:

△ start  
⊕ validate  
☆ finish



# Is Validation = 2PL?



S2:  $w_2(y) \quad w_1(x) \quad w_2(x)$

- Achievable with 2PL?
- Achievable with validation?

S2: w2(y) w1(x) w2(x)

- S2 can be achieved with 2PL:  
l2(y) w2(y) l1(x) w1(x) u1(x) l2(x) w2(x) u2(y) u2(x)
- S2 cannot be achieved by validation:  
The validation point of T2, val2 must occur before w2(y) since transactions do not write to the database until after validation. Because of the conflict on x, val1 < val2, so we must have something like  
S2: val1 val2 w2(y) w1(x) w2(x)  
With the validation protocol, the writes of T2 should not start until T1 is all done with its writes, which is not the case.







# Validation subset of 2PL?

- Possible proof (Check!):
  - Let  $S$  be validation schedule
  - For each  $T$  in  $S$  insert lock/unlocks, get  $S'$ :
    - At  $T$  start: request read locks for all of  $RS(T)$
    - At  $T$  validation: request write locks for  $WS(T)$ ; release read locks for read-only objects
    - At  $T$  end: release all write locks
  - Clearly transactions well-formed and 2PL
  - Must show  $S'$  is legal (next page)

- Say  $S'$  not legal (due to w-r conflict):  
 $S': \dots l1(x) \quad w2(x) \quad r1(x) \quad val1 \quad u1(x) \dots$ 
  - At val1: T2 not in Ignore(T1); T2 in VAL
  - T1 does not validate:  $WS(T2) \cap RS(T1) \neq \emptyset$
  - contradiction!
- Say  $S'$  not legal (due to w-w conflict):  
 $S': \dots val1 \quad l1(x) \quad w2(x) \quad w1(x) \quad u1(x) \dots$ 
  - Say T2 validates first (proof similar if T1 validates first)
  - At val1: T2 not in Ignore(T1); T2 in VAL
  - T1 does not validate:  
 $T2 \notin FIN \text{ AND } WS(T1) \cap WS(T2) \neq \emptyset$
  - contradiction!

# Conclusion:

## Validation subset 2PL



Validation (also called optimistic concurrency control) is useful in some cases:

- Conflicts rare
- System resources plentiful
- Have real time constraints

# Summary

Have studied C.C. mechanisms used in practice

- 2 PL
- Multiple granularity
- Tree (index) protocols
- Validation